Cavi Agentic Operations Platform (AOP)

Enterprise Capability Guide — B2B Offering Overview

> ***66+ Agents | 8 Divisions | 24/7 Execution | 4 Enterprise Tiers***
>
> *The roster expands continuously. These numbers represent a floor, not a ceiling.*

---

1. Overview — What Is the Cavi AOP?

The Cavi AOP is a managed, multi-agent AI workforce — coordinated AI teams that give humans superpowers.

Rather than selling a single chatbot or a prompt template library, Cavi deploys structured divisions of specialist agents — each with a defined role, explicit permissions, and an accountable chain of command. Every task has an owner. Every output has a validation gate. Every handoff is logged.

The platform runs 66+ agents across 8 divisions — and growing. Each division mirrors a real operational function (Engineering, Marketing, Research, Project Ops, and more), staffed by agents who know their domain and stay in their lane.

Deployment options:

- **Fully Managed** — Cavi provisions, monitors, and maintains your fleet end to end.
- **Hybrid** — Cavi agents operate alongside your existing teams and tools, integrated via API, Slack, Discord, or Microsoft Teams.
- **On-Premise / Private Cloud** — Kubernetes-native architecture deploys into your infrastructure with no data leaving your boundary.

> *The goal is not to replace your team. It is to give them a tireless, structured, always-on support layer that handles the operational weight so humans can focus on judgment, strategy, and creativity.*

---

2. Architecture at a Glance

```
                    ┌──────────────────────┐
                    │   Chief Orchestrator  │
                    │  (Executive Ops Hub)  │
                    └──────────┬───────────┘
                               │
        ┌──────────┬───────────┼──────────┬──────────┐
        │          │           │          │          │
┌───────▼──┐ ┌─────▼────┐ ┌────▼─────┐ ┌──▼───────┐ ┌─▼──────────┐
│Engineering│ │ Marketing │ │AI Media │ │ Talent   │ │Project Ops │
│& Dev (15) │ │& Brand(25)│ │Prod (5) │ │Acq. (5) │ │& Analytics │
└──────────┘ └──────────┘ └─────────┘ └─────────┘ │    (7)     │
                                                   └──────────┘

        ┌──────────┬──────────┐
        │          │          │
┌───────▼──┐ ┌─────▼────────┐ │
│ Research  │ │Infrastructure │ │
│& Intel(4) │ │& Reliability  │ │
└──────────┘ │     (4)       │ │
```

```
                  └──────────────┘ |
                                   |
                     (Central Orchestration
                      routes all divisions)
```

> *One routing authority. One dispatch layer. Every request enters through the Chief Orchestrator, gets classified, and lands in the right division with full context and a clear mandate.*

---

3. Division Details

Each division below operates as an autonomous unit — its own workspace, its own lead, its own specialist roster. Cross-division work is coordinated through the central orchestration layer with full audit trail.

3.1 Executive Operations — Central Hub

Field	Detail
Agents	1
Lead	Chief Orchestrator

Capabilities:

- Intake and triage of all inbound requests
- Cross-domain routing with ownership-explicit dispatch
- Operator-facing reporting and escalation handling
- Real-time session management and task status visibility
- Approved critical operations with full audit trail

Integrations: `Discord` · `Slack` · `Microsoft Teams` · `REST API` · `WebSocket RPC`

> *The Chief Orchestrator is the single point of entry and the single source of truth for task state. Nothing moves without a routing decision, and every routing decision is logged.*

3.2 Engineering & Development

Field	Detail
Agents	15
Lead	Engineering Orchestrator

Team: Frontend Architect · Backend Engineer · Infrastructure Engineer · CI/CD Engineer · Post-Deploy Validator · Git Operations Lead · Smart Contract Engineer · Database Engineer · Security Engineer · API Integration Specialist · Performance Engineer · Mobile Engineer · DevOps Engineer · Documentation Engineer

Capabilities:

- Frontend UI/UX development (React, TypeScript)
- Backend service and API development
- Smart contract development and auditing
- Kubernetes deployment and Helm chart management
- CI/CD pipeline management and optimization

- Post-deploy validation and smoke testing
- Git PR workflows and branch strategy enforcement

Integrations: `GitHub` · `GitHub Actions` · `Docker` · `Kubernetes` · `PostgreSQL` · `Redis` · `Foundry`

3.3 Marketing & Brand

Field	Detail
Agents	25
Lead	Marketing Orchestrator
Departments	5

Content & Copy: Story Architect · PR/Comms Lead · SEO Strategist · Product Marketing Lead

Growth & Acquisition: Growth Engineer · Paid Media Lead · Social Amplification Lead

Brand & Creative: Visual Alchemist · Creative Engine Lead · Brand Copywriter · Brand Copywriter II

Analytics & Measurement: Data Analytics Lead · Measurement Ops Lead · Trend Oracle

Community & Partnerships: Community Architect · Partnership Scout · Lifecycle CRM Lead · Stakeholder Ops Lead

Additional Specialists: Investor Deck Architect · Legal Risk Lead · Web3 Ecosystem Analyst · Token Messaging Lead · Event Ops Producer · Influencer Relations Lead

Capabilities:

- Campaign strategy and multi-channel execution
- Investor narrative and deck production
- Brand system and voice consistency enforcement
- SEO strategy and content calendar management
- Paid media setup, optimization, and reporting
- Community management and retention programs
- Partnership identification and activation workflows

Integrations: `Webflow` · `Notion` · `Canva API` · `Twitter/X` · `LinkedIn` · `Google Analytics`

3.4 AI Media Production

Field	Detail
Agents	5
Lead	AI Media Director

Team: Voice & Audio Specialist · Visual & Motion Specialist · Meme & Remix Specialist · Production Compliance Lead

Capabilities:

- AI voice synthesis and text-to-speech production
- Music production and audio branding
- AI image generation and visual asset creation
- Video, motion graphics, and animatics production

- Meme systems and remix template management
- Provenance tagging and rights documentation

Integrations: `ElevenLabs` · `Suno` · `Runway` · `Midjourney` · `DALL-E` · `FFMPEG`

3.5 Talent Acquisition

Field	Detail
Agents	5
Lead	Talent Operations Lead

Team: Opportunity Qualifier · Document Tailoring Specialist · Submission Automation Lead · Pipeline Manager

Capabilities:

- Job and candidate sourcing across platforms
- Opportunity scoring and qualification
- Company research and competitive dossiers
- Resume and cover letter tailoring per opportunity
- Application submission automation
- Pipeline state management and follow-up tracking

Integrations: `LinkedIn` · `Wellfound` · `Greenhouse` · `Lever` · `Playwright`

3.6 Project Operations & Analytics

Field	Detail
Agents	7
Lead	Project Ops Director

Team: Board/Kanban Manager · Deck & Presentation Specialist · Memo & Brief Writer · Comms Delivery Specialist · Data Visualization Lead · Handoff & Recap Specialist

Capabilities:

- GitHub Projects, Linear, and Jira board synchronization
- Blocker capture and escalation workflows
- Sprint standups and status report generation
- Deck and presentation production
- Memos, briefs, and executive summaries
- Chart and graph spec generation for data visualization

Integrations: `GitHub Projects` · `Linear` · `Jira` · `Slack` · `Notion` · `Google Workspace`

3.7 Research & Intelligence

Field	Detail
Agents	4
Lead	Research Director

Team: Internal Discovery Specialist · External Research Specialist · Synthesis Specialist

Capabilities:

- Web search and external evidence gathering
- Internal codebase and document discovery
- Competitive intelligence and market research
- Source-backed synthesis (evidence-first methodology)
- Technical reconnaissance and API documentation review
- Research packets for decks, briefs, and strategic planning

Integrations: `Web Search` · `Internal Docs` · `API Docs` · `Market Data`

3.8 Infrastructure & Reliability

Field	Detail
Agents	4
Lead	Infrastructure Lead

Team: Channel Integration Specialist · Session & Memory Ops Specialist · System Extensions Lead

Capabilities:

- Gateway runtime management and monitoring
- Cron job management and scheduling
- Discord bot operations and channel routing
- Plugin and connector lifecycle management
- Session and memory operations
- Log rotation and config drift monitoring

Integrations: `Kubernetes` · `Docker` · `Discord API` · `Slack API` · `HashiCorp Vault`

---

4. Enterprise Offerings

Cavi operates on a four-tier service model. Each tier builds on the previous, and all tiers include the accountability-first approach that defines the platform.

Tier 1 — Implementation

> *Get your fleet built, configured, and running.*

- Requirements discovery and division-fit mapping
- Agent identity and persona design
- Infrastructure provisioning (Kubernetes, cloud, or hybrid)
- Tool permission policy and RBAC configuration
- Gateway and channel integration setup
- Acceptance testing and handoff documentation

Tier 2 — Managed Service

> *Keep your fleet healthy, current, and performing.*

- 24/7 fleet health monitoring via automated daemons
- Config drift detection and remediation
- Memory and context lifecycle management
- Incident response SLA: gateway, sessions, routing
- Cron job and scheduled task management
- Model assignment updates as new models release
- Monthly platform health reports

Tier 3 — Custom Build

> *Extend the platform to fit your exact operational shape.*

- Custom specialist design: new agent roles, skill packs
- Domain-specific integrations (ERP, CRM, proprietary APIs)
- Custom model routing (open-source, private, third-party LLMs)
- Smart contract integration for on-chain workflows
- Branded AI identity system (personas, voices, tone guides)

Tier 4 — Upgrade & Expansion

> *Grow the fleet as your organization grows.*

- New division rollouts as your org scales
- Agent headcount scaling (sub-agent fleet expansion)
- Model upgrades as frontier models improve
- New channel integrations (platforms, inbound sources)
- Workflow automation additions (file-drop, webhooks, cron)

---

5. Target Sectors

Sector	Primary Use Cases
Technology & SaaS	Engineering orchestration, CI/CD operations, developer productivity, product marketing
Financial Services & FinTech	Reporting and analytics, investor comms, regulatory briefs, talent ops
Healthcare & Life Sciences	Research synthesis, clinical trial docs, compliance comms, talent sourcing
Media & Entertainment	AI media production, content ops, campaign execution, creator partnerships
eCommerce & Retail	SEO content, paid media, community ops, product launches
Web3 & Digital Assets	Smart contract dev, token messaging, NFT campaigns, ecosystem positioning
Professional Services	Client comms, project ops, deck production, brief and memo generation
Talent & Recruiting	End-to-end talent pipeline, candidate qualification, document tailoring
Startups & Scale-ups	Full-stack agentic ops from day one
Enterprise Conglomerates	Division-aligned deployments, cross-department routing, unified project ops

---

6. Token-Efficient Architecture

Efficiency is not an afterthought — it is built into the routing layer. The Cavi AOP is designed to minimize token overhead without sacrificing capability or context quality.

Bootstrap Deduplication

- **Orchestrators (8):** Full workspace root + domain skill pack. These agents carry the context needed to route, delegate, and validate.
- **Leaf Workers (58+):** Skill definition only (`skipBootstrap`). Each specialist receives only the task packet it needs — cutting per-task context overhead by 60-70%.

Model Tier Routing

Not every task needs a frontier model. The platform routes each task to the appropriate tier:

Task Type	Model Tier
Complex orchestration and cross-division routing	Frontier (Opus / Sonnet)
Specialist execution and domain work	Mid-tier
Health checks and status polling	Low-cost
Cron jobs and scheduled maintenance	Minimum viable

Zero-Token Infrastructure

Fleet health, config drift detection, log rotation, and memory archiving run as daemon scripts — no AI turns, no tokens consumed. Infrastructure does not eat into your operational budget.

Lean Routing

A single routing authority (`agents.yaml`) eliminates routing ambiguity. One file defines every agent, every permission boundary, and every dispatch path. No duplicate routing logic. No conflicting agent registries.

---

7. Foundation Architecture — PaperclipAI

Cavi is evaluating PaperclipAI as a foundational architecture — an open-source layer that provides core multi-agent orchestration primitives the platform can build upon.

PaperclipAI offers:

- **Agent lifecycle management** — Instantiation, teardown, and state management for agent processes
- **Message routing primitives** — The low-level dispatch and inter-agent communication protocols
- **Context window management** — Efficient token allocation and context threading across agent hierarchies
- **Tool execution framework** — Sandboxed tool invocation with permission boundaries

Cavi's roadmap calls for building on top of this foundation to deliver the operational layer: division structure, RBAC policy enforcement, operator-facing interfaces, fleet health monitoring, and the accountability model that makes the platform enterprise-ready.

> *PaperclipAI provides the engine primitives. Cavi is building the vehicle — purpose-built for enterprise operations, with the controls, safety systems, and instrumentation that production workloads demand.*

---

8. Data Ownership — Three-Layer Model

Data ownership in the Cavi AOP follows a clear, three-layer structure. There is no ambiguity about who owns what.

Layer	Scope	Owner
Layer 1 — Orchestration	Routing logic, agent definitions, dispatch rules, fleet configuration, platform infrastructure	**Cavi**
Layer 2 — Specialist Execution	Skill packs, tool integrations, model routing policies, agent personas, execution frameworks	**Cavi**
Layer 3 — Outputs	All deliverables, generated content, reports, code, media, research, and data produced by the fleet on your behalf	**Client**

> ***The work product is yours.*** *Cavi owns the platform and the operational machinery. Clients own every output the platform produces for them — full stop.*

This model ensures:

- **No vendor lock-in on outputs** — Your content, code, and data are portable from day one
- **Clear IP boundaries** — Platform IP stays with Cavi; operational outputs belong to the client
- **Audit-ready separation** — Every layer has explicit boundaries for compliance and legal review

---

9. Security & Compliance

Access Control

- **RBAC enforcement:** Every agent operates within an explicit permission boundary. No agent can exceed its defined scope.
- **Safe-bin enforcement:** Approved command allowlists per agent role. Agents execute only what they are permitted to execute.
- **Operator-only overrides:** Security-scope expansions require explicit human approval. No agent can escalate its own permissions.

Injection Defense

- Prompt injection detection built into the ingestion layer
- External content treated as untrusted by default
- All external-facing endpoints use strict method allowlists

Isolation

- **Workspace isolation:** Each division operates in its own workspace. Cross-division access requires explicit routing through the orchestration layer.
- **Context isolation:** Sub-agents receive only the task packet they need — no ambient context leakage.
- **Kubernetes namespace isolation:** Separate deployments per responsibility domain. No shared runtime between divisions.

Audit & Traceability

- Full audit trail for every cross-division handoff
- Every task has an owner, a validation plan, and a completion gate
- Gateway logs all session activity with per-agent attribution

Secrets Management

- HashiCorp Vault + External Secrets Operator for Kubernetes
- Host-level secrets never cross-contaminated with container or Kubernetes tiers
- No plaintext secrets in agent context or memory surfaces — `op://` references only
- GitHub token roles explicitly separated per agent scope

---

10. Deployment Topology

The Cavi AOP is Kubernetes-native. The deployment architecture is designed for reliability, isolation, and horizontal scaling.

```
┌──────────────────────────────────────────────────────────┐
│                   Kubernetes Cluster                     │
│                                                          │
│  ┌──────────────┐  ┌──────────────┐  ┌──────────────┐  │
│  │ Namespace:   │  │ Namespace:   │  │ Namespace:   │  │
│  │ exec-ops     │  │ engineering  │  │ marketing    │  │
│  │              │  │              │  │              │  │
│  │ Chief        │  │ Eng Orch +   │  │ Mktg Orch +  │  │
│  │ Orchestrator│  │ 14 agents    │  │ 24 agents    │  │
│  └──────────────┘  └──────────────┘  └──────────────┘  │
│                                                          │
│  ┌──────────────┐  ┌──────────────┐  ┌──────────────┐  │
│  │ Namespace:   │  │ Namespace:   │  │ Namespace:   │  │
│  │ ai-media     │  │ talent       │  │ project-ops  │  │
│  │              │  │              │  │              │  │
│  │ Media Dir +  │  │ Talent Ops +│  │ Proj Dir +   │  │
│  │ 4 agents     │  │ 4 agents     │  │ 6 agents     │  │
│  └──────────────┘  └──────────────┘  └──────────────┘  │
│                                                          │
│  ┌──────────────┐  ┌──────────────┐  ┌──────────────┐  │
│  │ Namespace:   │  │ Namespace:   │  │ Shared Svcs  │  │
│  │ research     │  │ infra        │  │              │  │
│  │              │  │              │  │ Vault        │  │
│  │ Research     │  │ Infra Lead +│  │ Secrets Ops  │  │
│  │ Dir + 3     │  │ 3 agents     │  │ Monitoring   │  │
│  └──────────────┘  └──────────────┘  └──────────────┘  │
│                                                          │
│  ┌──────────────────────────────────────────────┐       │
│  │              Gateway Layer                   │       │
│  │ Discord · Slack · Teams · REST · WebSocket   │       │
│  └──────────────────────────────────────────────┘       │
└──────────────────────────────────────────────────────────┘
```

Key characteristics:

- Each division runs in its own Kubernetes namespace with dedicated resource quotas
- Gateway layer handles all inbound channels and routes to the Chief Orchestrator
- Shared services (Vault, monitoring, log aggregation) are isolated from agent workloads
- Horizontal pod autoscaling available per division based on task queue depth

- Zero-downtime deployments via rolling updates and health-check gates

---

11. Get Started

The Cavi AOP is built for operators who want structured, accountable AI operations — not another chatbot.

> ***Book a consultation*** *to map your operational needs to the right divisions, agents, and service tier.*

Next steps:

1. **Discovery call** — We assess your operational landscape and identify high-impact deployment targets
2. **Division mapping** — Match your needs to Cavi divisions and define the initial agent roster
3. **Pilot deployment** — Stand up your fleet in a controlled environment with validation gates
4. **Production cutover** — Go live with full monitoring, SLA coverage, and operator dashboards

---

Cavi Agentic Operations Platform

66+ agents across 8 divisions — and growing.

Coordinated AI teams that give humans superpowers.

www.cavivaultagents.io