# CAVI COLLECTIVE

## Sales Enablement Playbook

**Version 2.0 • The Definitive Field Guide**

For founders, account executives, and partners selling Cavi.

**INTERNAL USE ONLY • CONFIDENTIAL**

# Table of Contents


| Part | Title | Page |
|---|---|---|
| **1** | **Know Cavi** | Core positioning, platform overview, output projections, pre-launch framing |
| **2** | **Know the Buyer** | Three segments (equal priority), personas, buying triggers |
| **3** | **The Sales Motion** | Velocity framework, role definitions, workflow assessment, handoff |
| **4** | **Discovery & Qualification** | Velocity questions, discovery bank, fit assessment |
| **5** | **Pricing Guidance** | 9-tier ladder, pilot pricing, model budgets, discount authority |
| **6** | **Objection Handling** | 4 hard objections (verbatim) + 12 common objections |
| **7** | **Competitive Intelligence** | Three-way frame, WorkflowAS, PaperclipAI foundation, category maps |
| **8** | **Sales Scripts** | SMB discovery (30 min), Solopreneur fast-close (15 min), Enterprise pilot (45 min) |
| **9** | **Outreach Templates** | LinkedIn, cold email, follow-up, re-engagement, founding customer |
| **10** | **Proof Points & Case Study** | Projections, internal case study, founding customer frame |
| **11** | **Technical Breakdown Bonus** | Glossaries, architecture, sample technical Q&A, escalation paths |


# Part 1. Know Cavi

## 1.1 The One-Sentence Pitch

> ***"Cavi gives operators coordinated AI teams that give humans superpowers."***

Note: It says **superpowers** — capability expansion. Not "AI replaces your team."

**Secondary value statement** (for cost-driven conversations):

> "*Coordinated AI specialists working under your team's direction — building capacity, freeing budget, and creating workflows that compound over time.*"

Both are approved. Lead with the primary in headlines, decks, and first impressions. Lead with the secondary when the prospect opens with cost concerns, budget constraints, or headcount pressure.

---

## 1.2 The 30-Second Story

> *Most teams are stuck: hire slowly and grow slowly, or hire fast and burn cash. Cavi breaks that trade-off. We give operators a coordinated AI workforce —* ***66+ specialists organized across 8 coordinated divisions — and growing*** *— that plugs into their business in three days, runs under their direction, and scales output without scaling headcount. Your team approves, the AI executes. You get the leverage of a 50-person org with the focus of a 10-person team.*

**Use this** when someone says "So what does Cavi do?" at a conference, on a call, or in a cold email opener. Thirty seconds. Pause. Let them ask the next question.

---

## 1.3 The Platform in Plain English

Cavi is the **Agentic Operations Platform — AOP.** Think of it as a managed AI workforce, not a tool.

What that means concretely:

- **66+ AI specialists — and growing.** Each specialist has a narrow job, a defined scope, and a name. They are not chatbots. They are not general-purpose assistants. Each one does one thing well, coordinates with the others, and delivers reviewable work. The roster expands continuously as we develop new capabilities.
- **8 coordinated divisions.** Executive Operations, Engineering & Development, Marketing & Brand, AI Media Production, Talent Acquisition, Project Operations & Analytics, Research & Intelligence, and Infrastructure & Reliability. Specialists inside a division coordinate with each other automatically. Divisions coordinate across boundaries when work requires it. Additional divisions are in development.
- **Labor as a Service (LAaS).** Clients pay a monthly subscription per specialist or per department. Each subscription includes a **monthly Model Budget** that covers typical operating volume. If a client exceeds it, additional model usage is passed through at our wholesale rate with **zero markup**. No per-seat math, no surprises, full transparency.
- **Human in the loop at every gate.** Your team reviews, approves, and directs. Cavi never ships work that a human did not green-light. Every output passes through an approval gate before it reaches production. The human is always in command.
- **Live in 3 days.** Not weeks. Not a six-month implementation. Three days from pilot kickoff to first specialist deployed and producing reviewable output.

> ***What "Live in 3 days" actually means***
>
> *The 3-day promise refers to your* ***first specialist deployed and producing reviewable output****. That means: onboarding call on Day 1, configuration and integration on Day 2, first live output reviewed by your team on Day 3. Complex multi-division deployments set realistic timelines during scoping — but even those start producing value within the first week. The point is speed: you see real work before you have time to second-guess the decision.*

**Division Summary**

| # | Division | Agents | Focus |
|---|---|---|---|
| 1 | Executive Operations | 1 | Central intake, cross-domain routing |
| 2 | Engineering & Development | 15 | Full-stack engineering orchestration |
| 3 | Marketing & Brand | 25 (5 departments) | End-to-end marketing execution |
| 4 | AI Media Production | 5 | Voice, audio, image, video, meme systems |
| 5 | Talent Acquisition | 5 | Sourcing, qualification, document production |
| 6 | Project Operations & Analytics | 7 | Project state management, reporting |
| 7 | Research & Intelligence | 4 | Evidence retrieval and synthesis |
| 8 | Infrastructure & Reliability | 4 | Runtime, gateway, integrations |

**Total: 66+ agents across 8 divisions — and growing.**

---

## 1.4 The Output Projections — Honestly

> ***IMPORTANT — READ THIS BEFORE EVERY CONVERSATION***
>
> *We do **NOT** have external client results. We are pre-launch. Every claim in this section must be framed as a projection from internal usage. If you present a projection as a customer case study, you are lying. Do not do it.*

**Projection 1: 8-10x output capacity gain**

> *"Based on Cavi's partner running the platform internally, we project an **8 to 10 times output capacity gain** when a coordinated AI team is paired with a human operator. We're sharing that as a working projection, not a customer case study, because we're early and we'd rather you trust us in two years than oversell you today."*

**Projection 2: 1 specialist handles workload of 1-3 employees**

> *"Based on Cavi's partner running the platform internally, we project a **single specialist can handle the workload of 1 to 3 employees** depending on the workflow. A specialist runs continuously, never takes vacation, and handles multi-step workflows in parallel. Same honesty framing — projection from internal usage, not a customer case study."*

**How to deliver projections in conversation:**

1. State the number clearly.
2. Immediately frame it as internal.
3. Acknowledge pre-launch status.
4. Pivot to the pilot: "That's why we offer the pilot — so you can see your own numbers."

---

## 1.5 Pre-Launch Is an Asset, Not a Liability

When a prospect asks "How many customers do you have?" — and they will — do not flinch. Lean in.

> *"We're pre-launch and that's deliberate. We're hand-selecting a small group of founding customers who get **founder-level attention**, **founding pricing protections**, and **a real seat at the table on what we build next**.*

> *Early customers shape the product. Later customers use the product. We're looking for people who want to shape it."*

**Why this works:**

- It reframes scarcity as exclusivity.
- It signals that the founders are in the room, not delegating to a support team.
- It creates urgency: the founding window closes.
- It attracts operators who want influence, not just software.

**If they push back:** "Fair question. We've been running the platform internally for months. The output projections come from real work, not demos. The pilot exists specifically so you can see it working on your workflows before you commit."

---

## 1.6 Things You Must Never Say

| Never say this | Say this instead |
|---|---|
| "We replace your team." | "We make your team superpowered." / "We give your team leverage." |
| "Our clients see..." | "We project..." / "Based on internal usage..." |
| "Set it and forget it." | "Your team reviews and approves everything." / "Human in the loop at every gate." |
| Any competitor name in writing or on a deck. | Use category language: "vertical tools," "traditional agencies," "general-purpose AI." |
| "Autonomous" (describing specialists) | "Coordinated," "managed," or "working under your team's direction." |
| Any compliance certification we do not have (especially SOC 2 Type II) | "We're building toward enterprise compliance. Here's what we do today." |
| "AI agents" in client-facing materials without context | "AI specialists" or "coordinated AI workforce." |

> ***Note on cost-cutting language:*** *"Save money," "cut costs," and "return on investment" are* ***approved phrases****. We do not shy away from economic value. Capability expansion and cost savings are equal-weight value drivers. Use whichever fits the conversation. If the prospect leads with budget pressure, lead with savings. If they lead with capacity pressure, lead with superpowers. Both are true.*

---



---

# Part 2. Know the Buyer — THREE SEGMENTS, EQUAL PRIORITY

---

> ***IMPORTANT: Cavi sells into three segments with EQUAL priority.*** *Each has a different buyer, a different pain pattern, and a different close. Learn all three. Do not blend them. Do not skip the solopreneur because the check is smaller. Do not skip the enterprise because the cycle is longer. Every segment matters. Every deal matters.*

---

## 2.1 Three Segments, Equal Priority

| Segment | Revenue Range | Decision Maker | Why This Segment Matters |
|---|---|---|---|
| **Solopreneur / 1-person ops** | $0-$2M | The founder (they ARE the company) | Highest emotional fit. They feel the pain of being alone. Solo tier at $1,200/mo is a no-brainer if discovery lands. Fastest close, highest satisfaction, strongest word-of-mouth. |
| **SMB (20-500 employees)** | $2M-$50M | Founder, COO, VP Marketing, VP Ops | Fastest decisions, real budgets, founders still in the room. Our pricing fits without a calculator. Our 3-day deploy story matches their urgency. Bread and butter. |
| **Enterprise ($10M+)** | $10M-$500M+ | CTO, VP Ops, VP Innovation, Chief AI Officer | Bigger checks, deeper relationships, expansion potential. Pursue when the entry point is a real operator, not a procurement lead with a 200-question RFP and no authority. |

**The rule:** No segment is "primary." No segment is "secondary." They are different plays in the same playbook. Run whichever play the field gives you.

---

## 2.2 Solopreneur Persona — "The Founder of One"

**Who they are:**

- Solo founders, indie operators, consultants billing $150-$500/hr
- Building a SaaS, agency, content business, or productized service — alone
- Already paying for 5-15 SaaS tools and comfortable with subscription spend in the $200-$2,000/month range
- Identifies as "resourceful" — proud of doing more with less
- Active on Twitter/X, LinkedIn, indie hacker communities, or niche Slack groups
- Has tried VA services, freelancers, and every AI tool on Product Hunt

**What keeps them up at night:**

- "Every hour I spend on ops is an hour I am not spending on revenue."
- "I cannot afford a full hire and I am tired of stitching tools together."
- "I want leverage, not more software."
- "I am one illness away from my business stopping."
- "I know AI can help me but I do not have time to figure out how."

**Their decision process:**

They decide alone. Fast. Usually within one conversation. The question is not "Can I afford this?" — it is "Does this feel like it will actually work, or is it another shiny thing that takes more time than it saves?" They are skeptical because they have been burned by tools that promised automation and delivered configuration.

**How to sell to solopreneurs:**

The Solo tier at **$1,200/month** is the entire pitch. Frame it as "a specialist who works 24/7, costs less than one freelance project per month, and lives inside your business in three days." Solopreneurs need a **10-minute demo**

**and a clear next step**. Move fast. Do not over-explain the platform architecture. Do not talk about divisions. Talk about the one specialist that will change their Tuesday.

**Approved opening line for solopreneurs:**

> "*What if you had a specialist that handled [their specific pain] 24/7, and it cost less than a single freelancer invoice?*"

---

## 2.3 SMB Primary Persona — "The Stretched Operator"

**Who they are:**

- Founder, COO, VP Marketing, or VP Operations at a 20-500 person company
- Revenue between $2M and $50M
- Has tried hiring agencies, freelancers, offshore teams — at least one engagement failed
- Already uses ChatGPT, Claude, or Gemini personally but does not have an AI "team"
- Runs on a mix of Slack, Notion, HubSpot, Jira, and duct tape
- Their team is good but undersized for the workload

**What keeps them up at night:**

- "I have more demand than capacity to deliver against."
- "Hiring takes 90 days and costs me $80K loaded — and the person might not work out."
- "My team is burning out and I can see it."
- "I keep seeing AI demos and I do not know which ones are real."
- "The board wants an AI strategy and I do not have one."

**Buying triggers:**

- Lost a key employee (urgency spike)
- New funding round or revenue milestone (budget unlocked)
- Failed agency engagement (trust gap — we fill it)
- Board asks "What is your AI strategy?" (political pressure)
- Quarterly planning season (budget allocation window)
- Competitive pressure — a competitor shipped something AI-powered

**How they decide:**

In conversations, not committees. They want to feel the person across from them gets their business. The price needs to fit without a calculator. And they need to believe they can have a working specialist live within a week. Three things: **trust the person, trust the price, trust the speed.**

**How to sell to SMB:**

Lead with the Free Workflow Assessment. Move to Starter Team ($2,800/mo for 3 specialists) or a department bundle. The close is the pilot: "Let us put one specialist on your biggest bottleneck. Three days. If it does not work, you walk away." SMB deals should close in 7-14 days. If you are past 14 days, something is stuck — find it and address it.

---

## 2.4 Enterprise Persona — "The AI-Curious Executive"

**Who they are:**

- CTO, VP Operations, VP Innovation, Chief AI Officer, Head of Transformation
- Company revenue $10M+, often $50M-$500M
- Has a budget line item for "AI initiatives" or "workforce transformation"

- Has likely run one or two failed AI pilots already (they know what bad looks like)
- Surrounded by vendors promising AI transformation; deeply skeptical
- Needs to show measurable results to a board or executive team

**What keeps them up at night:**

- "I have to show measurable AI ROI to the board this year."
- "My last AI pilot was a science project that never shipped."
- "Procurement and security review will kill anything that is not enterprise-grade."
- "I need something my team can adopt without a 6-month training program."
- "Every vendor says they are different. None of them are."

**How they decide:**

Through process. Evaluation, pilot, security review, procurement, legal. The cycle is longer — 30-90 days — but the contracts are bigger and stickier. The key: find the **operator** inside the enterprise. The person who actually does the work and feels the pain. Sell to them. Let them sell internally.

**How to sell to enterprise:**

**30-day, $10,000 paid pilot.** Not a free trial. Free trials attract tire-kickers. Paid pilots attract decision-makers. Lead with four pilot deliverables:

1. **Workflow Constraint Map** — Where their current process breaks down
2. **Risk & Control Design** — How Cavi handles their specific security and compliance requirements
3. **Implementation Blueprint** — Phased rollout plan with timelines and success metrics
4. **Executive Decision Ready Readout** — Board-ready summary of pilot findings, ROI projection, and go/no-go recommendation

> "*The pilot is not a demo with a timer. It is a 30-day engagement with four deliverables that your team can take to the board. At the end, you have the evidence to say yes or no. Either way, you walk away with the Workflow Constraint Map and the Implementation Blueprint.*"

---

## 2.5 Buying Triggers Cheat Sheet

| Trigger | Solo | SMB | Enterprise |
|---|---|---|---|
| Headcount loss / key employee departure | X | X | |
| New funding round or revenue milestone | | X | X |
| Failed agency or vendor engagement | | X | |
| Board asks "What is your AI strategy?" | | | X |
| Competitive pressure — competitor shipped AI feature | | X | X |
| Quarterly or annual planning season | X | X | X |
| Tool fatigue / SaaS sprawl | X | X | |
| Founder burnout / time poverty | X | | |
| Failed internal AI pilot | | | X |
| New executive hire (CTO, VP Ops, Chief AI Officer) | | | X |

| | | | |
|---|---|---|---|
| Seasonal demand spike approaching | | X | X |
| M&A integration — need to do more with combined teams | | | X |

**How to use this table:** When you identify a trigger in discovery, anchor every subsequent conversation to it. "You mentioned you lost your head of content last month. That is exactly the kind of gap Cavi fills — not as a replacement for that person, but as a force multiplier for whoever picks up the work."

---



---

# Part 3. The Sales Motion

---

## 3.1 The Velocity Framework

Cavi sells on velocity. Not pressure — velocity. The goal is to compress the time between "first conversation" and "working specialist deployed" to the absolute minimum, because the product sells itself once it is running.

**The 7-14 Day Deal Cycle:**

| Day | Milestone | Owner |
|---|---|---|
| **Day 0** | First conversation (inbound or outbound) | Account Executive or Founders |
| **Day 1-3** | Free Workflow Assessment delivered | Account Executive (Tier A/B) or Founders (Tier C) |
| **Day 4-7** | Decision conversation — pricing presented, objections handled | Account Executive + Founders (if enterprise) |
| **Day 7-14** | Contract signed and first specialist live | Account Executive owns close, delivery begins |

**The 14-Day Rule:** If a deal is not closing within 14 days of the first conversation, something is blocking it. Do not wait. Ask directly:

> "*We are at Day [X]. What is the one thing that needs to happen for you to say yes or no? I would rather hear no today than maybe next month.*"

This is not aggressive — it is respectful of their time and yours. Operators appreciate directness.

---

## 3.2 Who Sells What — Role Definitions

| Role | Owns | Does Not Own |
|---|---|---|
| **Founders** | Enterprise pilots, strategic accounts, escalations, founding customer relationships. Support Account Executives on calls when needed. Set pricing policy and discount authority. | Solo and Duo cycles after the first 10 deals close. Day-to-day SMB pipeline management. |

| | | |
|---|---|---|
| **Account Executives** | SMB primary motion, Starter Team and Department Bundle deals, Solo/Duo at scale. Own the pipeline, own the follow-up, own the close. | Enterprise security reviews, MSA negotiation, discount authority above 25%. |
| **Partners / Referral Channel** | Solopreneur and Solo tier deals, vertical-specific introductions, warm referrals. | Direct enterprise pursuit. Price negotiation beyond standard tiers. |

**Escalation rules:**

- If an Account Executive needs Founders on a call: 24-hour notice, brief in writing (who, what they need, what the block is).
- If a partner surfaces an enterprise lead: hand off to Founders within 48 hours with full context.
- If anyone encounters a security or compliance question they cannot answer: escalate to the engineering lead. Do not guess.

---

## 3.3 Free Workflow Assessment — Ownership Tiers

The Free Workflow Assessment is the top-of-funnel offer for every segment. It is free, it is valuable, and it is the fastest way to build trust and qualify a prospect simultaneously.

**What it is:** A 30-minute conversation (video or phone) that produces a **Mini Workflow Map** — a one-page visual showing the prospect's current workflow, where the bottlenecks are, and where Cavi specialists would plug in.

**What it is not:** A sales pitch disguised as a consultation. Do real work. Identify real bottlenecks. If Cavi is not the right fit, say so. The assessment builds trust whether or not it leads to a deal.

| Tier | Segment | Owner | Duration | Turnaround |
|---|---|---|---|---|
| **Tier A** | Solo and Solopreneur deals | Account Executives or Partners | 30 minutes | Immediate — Mini Workflow Map delivered same day |
| **Tier B** | SMB deals | Account Executives (Founders available for strategic SMB accounts) | 30-45 minutes | 48-hour turnaround on Mini Workflow Map |
| **Tier C** | Enterprise deals | Founders join the Account Executive on the call | 45-60 minutes | Evolves into paid 30-day Enterprise Pilot |

**After the assessment:**

- **Tier A:** "Based on what we mapped, here is the one specialist that would make the biggest difference. Want to try it for 14 days, no card required?"
- **Tier B:** "I mapped three bottlenecks. Here is how a Starter Team of three specialists would address all of them. Can we schedule a decision conversation for Thursday?"
- **Tier C:** "This is more complex than a single assessment can cover. The right next step is our 30-day Enterprise Pilot — $10,000, four deliverables, and your team walks away with a complete implementation blueprint whether you continue with Cavi or not."

---

## 3.4 Sales-to-Delivery Handoff

**Solo / Duo / Starter Team:** ONE PERSON owns end to end. The person who closes the deal is the person who onboards the client. No handoff. No "your account manager will reach out." The client sees one face from first conversation to first output. This is a feature, not a limitation.

**Department Bundle and Enterprise:** Hand off at contract signing. The Account Executive (with Founders supporting on enterprise deals) closes and writes a brief — one page, covering: client name, key contact, what they bought, what the first deliverable is, timeline, and any promises made during the sale. The delivery team picks up from there.

**The handoff brief must include:**

1. Client name and key contact (name, email, phone, preferred communication channel)
2. What they bought (tier, number of specialists, divisions activated)
3. What the first deliverable is (specific, not vague — "weekly content calendar for LinkedIn" not "marketing help")
4. Timeline commitments made during the sale
5. Any special terms, discounts, or promises
6. Known concerns or objections that were addressed (so delivery does not re-open them)

---

## 3.5 Flexible Fit

Cavi scales from a 2-agent bundle to a full multi-division deployment. If a prospect has workflows, we can add leverage. The platform adapts to the business — not the other way around.

**No business is too small.** A solopreneur with one specialist is a real client with real needs. Treat them accordingly.

**No business is too complex.** For compliance requirements still in progress, the pilot program is designed to build confidence and surface any gaps early. For unusual workflows, the Free Workflow Assessment maps the path before any commitment.

**The principle:** We do not turn prospects away because they do not fit a predefined box. We map their workflows, identify where leverage is possible, and propose the right configuration. If a prospect needs something we do not have yet, that is a roadmap conversation — not a disqualification.

> *"Cavi is not a one-size-fits-all product. It is a workforce that configures to your business. Tell us what you need done, and we will show you how the specialists handle it — or tell you honestly if we are not there yet."*

---



---

# Part 4. Discovery & Qualification

---

## 4.1 The Velocity Qualification Framework

Five questions. Every discovery call. Every segment. If you cannot answer all five by the end of the call, you do not have a qualified lead — you have a conversation.

| # | Category | The Question | What You Are Really Asking |
|---|---|---|---|
| 1 | **PAIN** | "What is the work that is not getting done right now?" | Is there a real gap, or are they browsing? |

| | | | |
|---|---|---|---|
| 2 | **URGENCY** | "What changes if you do not solve this in 30 days?" | Is there a deadline, a consequence, a trigger? |
| 3 | **AUTHORITY** | "Who else needs to nod for this to move?" | Are you talking to the decision maker, or do you need to get to them? |
| 4 | **BUDGET FIT** | "Does $1,200 to $7,500 a month fit your operating budget?" | Can they actually pay for this? |
| 5 | **FIRST DELIVERABLE** | "If we made one thing live in 3 days, what would it be?" | Can they articulate what they need? If yes, they are qualified. |

**How to use the framework:**

- **All five answered = qualified.** Move to pricing and close.
- **Four of five answered = almost qualified.** Identify the missing one and address it directly.
- **Three or fewer = not ready.** Offer the Free Workflow Assessment and re-engage when they have clarity.

---

## 4.2 Discovery Question Bank

### Opening Questions (First 5 minutes)

> "*Tell me about your business and what your team looks like today.*"

> "*What made you take this call? What prompted you to look at something like Cavi right now?*"

> "*How are you handling [the specific workflow they mentioned] today?*"

> "*Walk me through a typical week. Where do you spend time that you wish you did not?*"

> "*What have you tried before to solve this? What worked and what did not?*"

### Middle Questions (Minutes 5-20)

> "*If I could wave a wand and fix one thing in your operations tomorrow, what would it be?*"

> "*How many people touch that workflow today? What does it cost you — in time, money, or missed opportunities?*"

> "*What does your tech stack look like? What tools are you already using?*"

> "*Who on your team would interact with a Cavi specialist day to day?*"

> "*Have you used AI tools before? ChatGPT, Claude, anything else? What was that experience like?*"

> "*What does success look like in 90 days? If Cavi works perfectly, what is different about your business?*"

> "*Is there a specific event, deadline, or milestone driving the timeline here?*"

### Closing Questions (Last 5-10 minutes)

> "*Based on what you have told me, here is what I would recommend. Does that feel right to you?*"

> "*What would need to be true for you to say yes this week?*"

> "*Is there anything that would make you say no right now? I would rather hear it.*"

> "*Who else on your team should see this before you decide?*"

> *"Can we schedule the next conversation for [specific day and time]?"*

---

## 4.3 Velocity Signals — Read the Room

| Signal | What It Means | What to Do |
|---|---|---|
| They ask about pricing early | They are ready to buy. Do not delay. | Give the price. Do not make them wait for a "pricing call." |
| They ask about implementation timeline | They are visualizing using it. | Reinforce "live in 3 days" and offer to start immediately. |
| They mention a specific workflow or bottleneck | They have a real pain point, not abstract interest. | Anchor everything to that specific workflow for the rest of the call. |
| They say "I need to think about it" | There is an unspoken objection. | Ask: "What specifically do you need to think through? I might be able to answer it right now." |
| They ask "Can I talk to a current customer?" | Pre-launch concern. Normal. | Use the founding customer frame from Section 1.5 and offer the pilot. |
| They compare you to a specific tool or vendor | They are evaluating. They need differentiation. | Use the three-way frame from Part 7. Never name the competitor. |
| They start talking about their team's reaction | They are already thinking about adoption. | Pivot to the "superpowers" frame: "Your team gets amplified, not replaced." |
| They go silent after pricing | Sticker shock or internal calculation. | Wait. Let them process. Then ask: "What are you comparing that number to?" |
| They ask for a free trial | They want proof before commitment. | Solo/Duo/Starter: "14 days free, no card required." Larger: "Free Assessment first, then $1,500 pilot." |
| They forward your email to someone else | They are building internal consensus. | Ask who they forwarded to and offer to talk to that person directly. |

---

## 4.4 Fit Assessment

Every prospect gets a fair assessment. The goal is not to disqualify — it is to find the right fit and set honest expectations.

**For compliance requirements (SOC 2, HIPAA, FedRAMP):**

> *"We're building toward enterprise compliance. The pilot program is designed to surface requirements early and build confidence. Let us show you how we handle security today — encryption at rest and in transit, role-based access, full audit trails, and data isolation. For specific certifications on your roadmap, our engineering lead can walk through our compliance timeline."*

**For on-premise requirements:**

> *"We offer hybrid deployment options — Kubernetes on your infrastructure or Cavi-hosted managed service. Let's map the right topology during scoping. The Enterprise Pilot is designed to evaluate exactly this: what runs where,*

> *what data stays where, and how the architecture fits your requirements."*

**For complex procurement processes:**

> *"The 30-day Enterprise Pilot is built specifically to give your procurement and security teams the evidence they need — four deliverables, fixed scope, clear outcomes. We have been through evaluation processes before and we know what your team needs to see."*

**For RFP processes:**

> *"We participate in structured evaluations. The pilot deliverables — Workflow Constraint Map, Risk & Control Design, Implementation Blueprint, and Executive Readout — are built for exactly this purpose. They give your evaluation team concrete evidence, not marketing materials."*

**For prospects with very early-stage needs:**

> *"It sounds like you are still figuring out where AI fits in your business. That is exactly what the Free Workflow Assessment is for. Thirty minutes, no commitment, and you walk away with a workflow map whether you buy Cavi or not. Want to schedule that?"*

---



---

# Part 5. Pricing Guidance

---

## 5.1 The 9-Tier Ladder

| Tier | Specialists | Monthly Price | Per Specialist | Notes |
|---|---|---|---|---|
| **Free Workflow Assessment** | — | $0 | — | 30-minute call, no commitment, Mini Workflow Map delivered |
| **Solo** | 1 | $1,200 | $1,200 | Entry point for solopreneurs and small tests |
| **Duo** | 2 | $2,200 | $1,100 | Save $200/mo vs. 2x Solo |
| **Starter Team** | 3 | $2,800 | $933 | Save $800/mo vs. 3x Solo — best value entry |
| **Research & Intel** | 4+ | $2,400 | $600 | Full Research & Intelligence division |
| **Talent Acquisition** | 5+ | $2,800 | $560 | Full Talent Acquisition division |
| **AI Media Production** | 5+ | $2,800 | $560 | Full AI Media Production division |
| **Project Operations** | 7+ | $3,400 | $485 | Full Project Operations & Analytics division |
| **Engineering & Dev** | 15+ | $5,800 | $385 | Full Engineering & Development division |

| | | | | |
|---|---|---|---|---|
| **Marketing & Brand** | 25+ | $7,500 | $300 | Full Marketing & Brand division (5 departments) |

**Pricing logic to remember:** The per-specialist cost drops as you go up the ladder. Solo is $1,200/specialist. Marketing & Brand is $300/specialist. The more you buy, the cheaper each specialist gets. This is intentional — it rewards commitment and makes expansion natural.

---

## 5.2 Universal Rules

These rules apply to every tier, every deal, every segment. No exceptions.

- **$0 setup fees on every tier.** No implementation charges. No onboarding fees. No "professional services" line item. The subscription is the price.
- **Solo, Duo, Starter Team: 14 days free trial, no card required.** The prospect signs up, gets a specialist deployed, and uses it for 14 days. If they do not convert, they walk away. No friction. No "forgot to cancel" charges.
- **Live in 3 days.** First specialist deployed and producing reviewable output within 3 days of kickoff.
- **Month-to-month by default.** No annual contracts required. No long-term commitments. Cancel with 30 days notice.
- **Annual option available:** 2 months free (pay for 10, get 12). Offer this when the prospect signals long-term intent.

**Multi-Department Discount:**

| Departments | Discount |
|---|---|
| 2 departments | 15% off combined monthly |
| 3 departments | 20% off combined monthly |
| 4+ departments | 25% off combined monthly |

**Example:** A client subscribing to Marketing & Brand ($7,500) + Engineering & Dev ($5,800) = $13,300/mo before discount. With 2-department discount (15%): **$11,305/mo.** That is 40+ specialists across two divisions for less than the loaded cost of two mid-level employees.

---

## 5.3 Pilot Pricing

| Pilot Type | Price | Duration | What They Get |
|---|---|---|---|
| **SMB Mini-Pilot** | $1,500 | 3 days | Live specialist deployed on their workflow + Mini Workflow Map + 4 deliverables (Workflow Constraint Map, Risk & Control Design, Implementation Blueprint, Executive Readout) |
| **Enterprise Pilot** | $10,000 | 30 days | Full Decision Ready package: all 4 deliverables at enterprise depth, multi-specialist deployment, security review included, board-ready summary |

**When to offer each:**

- **SMB Mini-Pilot ($1,500):** When the prospect is qualified but wants to "see it work" before committing to a monthly subscription. The $1,500 is credited toward the first month if they convert.

- **Enterprise Pilot ($10,000):** When the prospect has enterprise requirements (security review, procurement process, board approval) and needs structured evidence to make a decision. The $10,000 is credited toward the first quarter if they convert.

**The pilot close:**

> *"Here is what I would recommend. Let us run a 3-day pilot on [their specific workflow]. You invest $1,500. We deploy a specialist, produce real output, and deliver four documents that map exactly how Cavi fits your business. If you move forward, the $1,500 is credited to your first month. If you do not, you keep the Workflow Map and the Blueprint. Either way, you walk away with something valuable."*

---

## 5.4 How to Talk About Price

**The per-specialist story:**

> *"Each Cavi specialist costs between $300 and $1,200 per month depending on the tier. A single specialist handles the workload we project at 1 to 3 employees — running 24/7, no vacation, no ramp-up time. Compare that to one employee at $80,000 to $120,000 per year loaded, plus 90 days to hire and 30 days to ramp."*

**The anchoring move:**

> *"What does it cost you to hire one person for this role? $80K? $100K? That is $6,700 to $10,000 per month loaded. Our Starter Team — three specialists — is $2,800 per month. For less than a third of the cost of one hire, you get three specialists that run around the clock."*

**The capacity multiplier:**

> *"We project that a team of 66+ specialists across 8 divisions delivers the capacity of a 50-person organization — based on internal usage. The Marketing & Brand division alone — 25 specialists at $7,500 per month — projects to the capacity of 8 to 15 full-time marketers. That is $7,500 versus $640,000 to $1.2 million per year in salary alone."*

**The pilot close on price:**

> *"If the monthly number feels like a leap of faith, start with the pilot. $1,500 for 3 days, real output, and the four deliverables. Let the work speak for itself."*

---

## 5.5 Discount Authority

| Role | Discount Authority | Approval Required | Notes |
|---|---|---|---|
| **Account Executive** | Up to 10% off list | No additional approval needed | Standard discount for competitive situations or fast close |
| **Account Executive** | Up to 25% off list | No additional approval needed | Strategic deals, founding customers, multi-department bundles |
| **Founders** | Above 25% | Requires written rationale and joint sign-off from both Founders | Reserved for exceptional circumstances — anchor partnerships, case study clients, first-in-vertical deals |

**Discount principles:**

1. **Never lead with a discount.** Present full price first. Always.

2. **Discounts are earned, not expected.** Annual commitment, multi-department bundle, founding customer status, case study participation — these earn discounts.
3. **Founding customer pricing is time-limited.** The founding window closes. Once it does, these prices are gone. Make that clear.
4. **Document every discount.** Written record in the CRM: who approved, why, and what the client committed to in return.

---

## 5.6 Model Budget and Wholesale Pass-Through

Every Cavi subscription includes a **monthly Model Budget** — a pre-allocated amount that covers the AI model usage for typical operating volume at that tier. The Model Budget is included in the subscription price. The client does not pay extra for normal usage.

**If a client exceeds their Model Budget:**

- Additional model usage is passed through at Cavi's wholesale rate.
- **Zero markup.** We pay the model providers X, we charge the client X. No hidden margin on model costs.
- Usage is tracked in real time and visible to the client in their dashboard.
- The client is notified before they exceed the budget, not after.

**Why this matters in the sale:**

> "*Every subscription includes a Model Budget that covers normal operating volume. If you run heavy one month and exceed it, the overage is passed through at our wholesale rate — zero markup. You see exactly what the models cost us, and you pay the same. No hidden fees, no surprise line items. Full transparency.*"

**Model Budget by Tier:**

| Tier | Subscription | Included Model Budget | Pass-Through Above |
|---|---|---|---|
| Solo | $1,200 | $150 | Wholesale, zero markup |
| Duo | $2,200 | $275 | Wholesale, zero markup |
| Starter Team | $2,800 | $375 | Wholesale, zero markup |
| Research & Intel | $2,400 | $500 | Wholesale, zero markup |
| Talent Acquisition | $2,800 | $600 | Wholesale, zero markup |
| AI Media Production | $2,800 | $1,200 | Wholesale, zero markup |
| Project Operations | $3,400 | $850 | Wholesale, zero markup |
| Engineering & Dev | $5,800 | $1,800 | Wholesale, zero markup |
| Marketing & Brand | $7,500 | $2,500 | Wholesale, zero markup |

**Note:** AI Media Production has a higher Model Budget relative to its subscription price because media generation (images, video, audio) consumes significantly more model resources than text-based workflows. This is factored into the pricing — no surprise for the client.

**The transparency pitch:**

> "*Most AI platforms bury model costs in their pricing and mark them up 3-5x. We show you exactly what the models cost and pass them through at wholesale. The subscription covers the platform, the specialists, the orchestration,*

> *and a generous Model Budget. If you go over, you pay what we pay. That is it."*

---



---

# Part 6. Objection Handling

---

> ***How to use this section:*** *The first four objections are the hard ones. Memorize them verbatim. Practice them until the words come out naturally. The remaining twelve are common objections you will hear regularly — know the framework, adapt the language to the conversation.*

---

## 6.1 The Four Hard Objections (Memorize Verbatim)

---

### Objection #1: "What about hallucination?"

**What they are really asking:** "Can I trust the output? Will it make things up and embarrass my company?"

**The response (verbatim):**

> *"Valid concern. Here is how we handle it — four layers of defense.*
>
> ***First: narrow scope.*** *Each specialist has a defined job and a defined scope. It is not a general-purpose chatbot trying to be everything. A specialist that writes product descriptions does not suddenly start writing legal contracts. Narrow scope means fewer opportunities to hallucinate.*
>
> ***Second: human approval gates.*** *Every output passes through your team for review before it reaches production. Nothing ships without a human green-light. This is not optional — it is built into the platform.*
>
> ***Third: full audit trails.*** *Every input, every output, every decision the specialist makes is logged and reviewable. If something goes wrong, you can trace exactly what happened and why.*
>
> ***Fourth: model routing with fallbacks.*** *We route each task to the most appropriate model for that type of work, with automatic fallbacks if the primary model is unavailable or underperforming. This means we are always using the best available model for each specific task, not forcing everything through one model.*
>
> *The honest version: hallucination risk is never zero with any AI system. But narrow scope plus human review plus audit trails reduces it to a level that is manageable and traceable. That is the standard we hold ourselves to."*

---

### Objection #2: "What if Cavi goes away?"

**What they are really asking:** "Am I building a dependency on a startup that might not exist in two years?"

**The response (verbatim):**

> *"Three layers of protection.*
>
> ***First: your data is exportable on demand.*** *Everything Cavi produces for you — every document, every output, every workflow — is yours and exportable at any time. You are never locked in because you always have your work.*
>
> ***Second: enterprise-grade cloud infrastructure.*** *We run on enterprise cloud infrastructure with redundancy, automated backups, and disaster recovery. The platform is built to stay up.*

> ***Third: contractual transition assistance.*** *Our MSA includes transition assistance provisions. If Cavi ever ceases operations, we are contractually obligated to provide your data, documentation, and reasonable support to help you transition.*
>
> *But the bigger answer is this: we are building a business, not running an experiment. We are funded, we are growing, and our founding customers get the most attention and the best pricing because we are building something we intend to run for a long time."*

---

### Objection #3: "Who owns the work product?"

**What they are really asking:** "If Cavi builds something for my business, do I own it? Or does Cavi?"

**The response (verbatim):**

> *"Three layers, clearly defined.*
>
> ***Layer 1 — Orchestration: Cavi owns.*** *The routing logic, intake triage, cross-domain coordination, and dispatch patterns. This is the nervous system of the platform — how work gets from your request to the right specialist and back.*
>
> ***Layer 2 — Specialist Execution: Cavi owns.*** *The agent skill packs, prompt engineering, model routing logic, and execution frameworks. These are the capabilities that make each specialist effective at their job.*
>
> ***Layer 3 — Outputs: You own.*** *Everything we build for you. Your data, your inputs, your finished work products, and any business deliverables that come out of the system. Code, content, designs, reports, strategies — all yours.*
>
> ***The simple version: you own what we build for you. We own how we build it.***
>
> *There is no overlap in the middle. Your work product is yours from the moment it is created. Our platform technology is ours. This is the same model as every professional services firm and every SaaS platform — you own the output, they own the tooling."*

---

### Objection #4: "OpenAI / Anthropic / Google will just build this."

**What they are really asking:** "Why should I bet on Cavi when the big AI companies might offer the same thing?"

**The response (verbatim):**

> *"Three reasons that will not change.*
>
> ***First: different layer.*** *OpenAI, Anthropic, and Google build foundation models — the raw intelligence. Cavi builds on top of those models to create a managed workforce. They are the engine manufacturers. We are the fleet operator. Intel makes chips. Apple makes computers. Different layer, different business.*
>
> ***Second: model-agnostic.*** *We are not tied to any single model provider. We route between multiple providers — selecting the best model for each specific task. If OpenAI releases a better model tomorrow, our specialists get better tomorrow. If Anthropic leaps ahead on reasoning, we route reasoning tasks to Anthropic. The big AI companies compete with each other. We benefit from all of them.*
>
> ***Third: we sell to operators, not developers.*** *The big AI companies sell APIs, SDKs, and developer tools. They are building for engineers. We are building for the COO who needs 10,000 product descriptions by Friday and does not want to learn Python. Different buyer, different product, different market."*

---

## 6.2 The Twelve Common Objections

---

### Objection #5: "It is too expensive."

**What they are really asking:** "I do not see the value relative to the price — yet."

> *"What are you comparing it to? Let me run the math with you.*
>
> *A single employee costs $80,000 to $120,000 per year loaded — that is $6,700 to $10,000 per month. Plus 90 days to hire and 30 days to ramp. Our Starter Team — three specialists — is $2,800 per month. Deployed in 3 days.*
>
> *We project each specialist handles the workload of 1 to 3 employees based on internal usage. So $2,800 per month for the projected capacity of 3 to 9 employees. Compare that to the cost of even one hire.*
>
> *But if the monthly number is a stretch right now, start with Solo at $1,200. Or start with the Free Workflow Assessment and let us show you where the leverage is before you spend anything."*

---

### Objection #6: "We already use ChatGPT / Claude / Gemini."

**What they are really asking:** "How is this different from what I already have?"

> *"That is a great starting point — it means your team already believes AI can help. The difference is tools versus workforce.*
>
> *ChatGPT is a tool. You open it, you type a prompt, you get a response, you close it. It does not know your business, it does not remember last week, and it does not coordinate with anything else.*
>
> *Cavi is a workforce. 66+ specialists organized into 8 coordinated divisions. Each specialist has a defined job, knows your workflows, coordinates with other specialists, and produces reviewable output on a schedule. It is the difference between having a search engine and having a research team.*
>
> *Keep using ChatGPT for ad hoc questions. Use Cavi for the work that needs to get done consistently, at scale, with coordination."*

---

### Objection #7: "I need to think about it."

**What they are really asking:** Something specific that they have not voiced yet.

> *"Totally fair. Can I ask — what specifically do you need to think through? Is it the price, the timing, the fit, or something else? I would rather address it now than have you spend a week wondering about something I could answer in two minutes."*

**If they cannot articulate it:** Offer the Free Workflow Assessment as a low-commitment next step.

> *"Here is what I would suggest — let us do the Free Workflow Assessment. Thirty minutes, no commitment. It will give you something concrete to think about instead of abstract questions. Can we schedule that for tomorrow?"*

---

### Objection #8: "We tried AI before and it did not work."

**What they are really asking:** "Convince me this is different from the last AI thing that wasted my time and money."

> *"Tell me about that experience. What did you try and where did it break down?"*

*Listen carefully. The answer tells you exactly what to address.*

> *"Most failed AI projects fall into one of two categories: science projects or tool sprawl. Science projects are when someone buys an AI platform, spends six months configuring it, and never gets to production output. Tool sprawl is when the team adopts five different AI tools that do not talk to each other and create more work than they save.*

> *Cavi is neither. It is a workforce deployment. Specialists are deployed on your workflows in 3 days, producing reviewable output immediately. No six-month configuration. No disconnected tools. One platform, coordinated specialists, human approval at every gate.*
>
> *The pilot exists specifically for this reason — so you can see it working on your workflows before you commit. Three days, real output. If it does not work, you walk away."*

---

### Objection #9: "We are too small for this."

**What they are really asking:** "Is this built for companies like mine, or am I going to be underserved?"

> *"Actually, smaller teams get MORE out of Cavi per dollar. Here is why.*
>
> *A 500-person company adding 3 Cavi specialists increases their capacity by less than 1%. A 5-person company adding 3 Cavi specialists increases their capacity by 60% or more — based on our projections.*
>
> *The Solo tier exists specifically for this. $1,200 per month, one specialist, deployed in 3 days. That is less than a single freelance project per month, and the specialist works 24/7. Our smallest clients often see the biggest relative impact because every hour of leverage matters more when you are running lean.*
>
> *Start with one specialist on your biggest bottleneck. If it works, expand. If it does not, you are out $1,200 and 3 days."*

---

### Objection #10: "We are too big / we need enterprise features."

**What they are really asking:** "Can your startup actually serve a company like ours?"

> *"We built for this. Here is what we have today:*
>
> - ***SSO integration** — your team logs in with your existing identity provider*
> - ***Full audit trails** — every input, output, and decision logged and exportable*
> - ***Role-based access control (RBAC)** — your team sees only what they should see*
> - ***MSA terms** — we do enterprise contracts with enterprise terms*
> - ***Data isolation** — your data does not commingle with other clients*
> - ***Hybrid deployment** — Kubernetes on your infrastructure or Cavi-hosted managed service*
>
> *For deeper technical questions — data residency, encryption standards, penetration testing, compliance roadmap — our engineering lead can walk through the details with your security team. We have done this before and we know what your team needs to see.*
>
> *The right next step is the 30-day Enterprise Pilot. $10,000, four deliverables, and your security team gets full access to evaluate the platform architecture."*

---

### Objection #11: "My team will resist this."

**What they are really asking:** "Will my people see this as a threat to their jobs?"

> *"This is the most important conversation to get right. Here is how we frame it — and how we recommend you frame it to your team:*
>
> ***Cavi gives your team superpowers. They are not being replaced — they are being amplified.***
>
> *The specialists handle the repetitive, time-consuming, high-volume work that burns your team out. Your people focus on the strategic, creative, high-judgment work that only humans can do.*

> *Think about it this way: when your team got access to email, they did not lose their jobs — they got faster. When they got project management tools, they did not lose their jobs — they got more organized. Cavi is the same pattern: more leverage, not less people.*
>
> *The teams that adopt Cavi fastest are the ones that are already overwhelmed. When someone is drowning in work and you hand them a life raft, they do not complain about the color of the raft.*
>
> *We recommend announcing it as: 'We are investing in AI support so you can focus on the work you are best at and stop drowning in the work you hate.'"*

---

### Objection #12: "Show me the ROI math."

**What they are really asking:** "Give me a number I can put in a spreadsheet for my boss."

**Conservative version (labor offset):**

> *"Conservative math: assume one specialist saves your team 10 hours per week of manual work. At a blended rate of $50 per hour, that is $2,000 per month in labor offset. The Solo tier is $1,200 per month. Net savings: $800 per month minimum, plus your team's 10 hours per week redirected to higher-value work.*
>
> *Scale that to a Starter Team — 3 specialists — and you are looking at 30 hours per week of labor offset. $6,000 per month in value against $2,800 per month in cost. That is a 2.1x return before you factor in the compounding effect of specialists running 24/7."*

**Capacity multiplier version:**

> *"The capacity multiplier math: we project each specialist handles the workload of 1 to 3 employees. A Starter Team of 3 specialists at $2,800 per month projects to the capacity of 3 to 9 employees. At $80,000 per year loaded per employee, that is $240,000 to $720,000 per year in equivalent capacity — for $33,600 per year in Cavi subscriptions. That is a 7x to 21x capacity return.*
>
> *We frame this as a projection from internal usage, not a guarantee. But the pilot exists specifically to generate your own numbers."*

---

### Objection #13: "What if I want to leave?"

**What they are really asking:** "Am I going to be locked in?"

> *"Month-to-month. No penalty. 30 days notice. Your data is exportable at any time.*
>
> *We do not lock you in because we do not need to. If Cavi is working, you will stay because it is working, not because a contract forces you to. If it stops working, we would rather know about it and fix it than hold you hostage.*
>
> *The annual option — pay for 10 months, get 12 — is available if you want to save money after you have seen it work. But we never push annual contracts upfront. Earn the commitment, then ask for it."*

---

### Objection #14: "How is this different from an agency?"

**What they are really asking:** "I have used agencies before. Why would I switch to AI specialists?"

> *"Four differences that matter:*
>
> ***Speed.*** *An agency takes 2-4 weeks to onboard and start producing. Cavi deploys in 3 days. When you need something done, you need it done now, not after a kickoff meeting, a strategy session, and a creative brief review.*

> ***Scope.*** *An agency is typically good at one thing — design, or content, or development. Cavi gives you 66+ specialists across 8+ coordinated divisions — marketing, engineering, media production, research, project operations, and more. One platform, cross-functional coordination.*
>
> ***Incentives.*** *An agency is incentivized to expand the scope of work. Their revenue goes up when your project gets bigger. Cavi is a flat subscription. Our incentive is to make you successful and keep you subscribed. We win when you win, not when the project scope creeps.*
>
> ***Knowledge retention.*** *When you leave an agency, the institutional knowledge walks out the door. With Cavi, the specialists retain your workflow context, preferences, and standards. Nothing gets lost between engagements.*
>
> *The honest version: agencies are great for high-strategy, high-creativity work where you need a senior human brain. Cavi is great for high-volume, high-consistency execution where you need reliable output at scale. Some clients use both."*

---

### Objection #15: "Can I get a free trial?"

**What they are really asking:** "I want to test it before I pay."

> ***For Solo, Duo, or Starter Team:***
>
> *"Yes. 14 days free, no card required. You get a specialist deployed on your workflow, and you use it for two weeks. If it works, you subscribe. If it does not, you walk away. No friction, no 'forgot to cancel' charges."*

> ***For Department Bundles or Enterprise:***
>
> *"For larger deployments, the right starting point is the Free Workflow Assessment — 30 minutes, we map your workflows and identify where Cavi adds the most leverage. From there, we move to the SMB Mini-Pilot at $1,500 for 3 days or the Enterprise Pilot at $10,000 for 30 days. Paid pilots are not about the money — they are about ensuring both sides are invested in getting real results."*

---

### Objection #16: "What about hidden fees? Token costs?"

**What they are really asking:** "Is the price really the price, or will my bill balloon?"

> *"The subscription price is the price. Here is exactly how it works.*
>
> *Every tier includes a monthly Model Budget — a pre-allocated amount that covers AI model usage for typical operating volume. For example, the Starter Team at $2,800 per month includes $375 in Model Budget. For most clients, the included budget is more than enough.*
>
> *If you run heavy one month and exceed the Model Budget, the overage is passed through at our wholesale rate — zero markup. We pay the model provider X, you pay X. No hidden margin.*
>
> *You see your usage in real time in the dashboard. We notify you before you exceed the budget, not after. And there is never a surprise bill — you always know where you stand.*
>
> *There are no setup fees, no onboarding fees, no implementation fees, no 'professional services' charges. The subscription is the price. The Model Budget is included. Overages are at cost. That is it."*

---



---

# Part 7. Competitive Intelligence

---

> ***INTERNAL ONLY.*** *Never name a competitor in client-facing materials, decks, emails, or proposals. Use category language: "vertical tools," "traditional agencies," "general-purpose AI." This section is for your preparation only.*

---

## 7.1 The Public Three-Way Frame

When a prospect is evaluating alternatives, use this framework. It names categories, not companies. It is safe for any conversation.

| Option | What It Is | Where It Falls Short |
|---|---|---|
| **Build It Yourself** | Open-source frameworks, internal AI team, custom development | Months of build time, ongoing engineering cost, no cross-functional coordination, your engineers maintain AI instead of building product |
| **Buy a Vertical Tool** | Single-purpose AI tools (AI writer, AI designer, AI coder) | Lock-in to one function, no cross-functional coordination, tool sprawl, each tool has its own learning curve and billing |
| **Hire a Coordinated Workforce (Cavi)** | Managed AI workforce — 66+ specialists across 8 coordinated divisions, human in the loop, live in 3 days | — |

**How to use the frame:**

> *"When companies look at AI, they usually see three options. Build it themselves, which takes months and engineering resources. Buy point solutions — an AI writer here, an AI coder there — which creates tool sprawl. Or hire a coordinated workforce that covers multiple functions and works under your team's direction. We are the third option."*

---

## 7.2 WorkflowAS — Verified Internal Profile

> ***INTERNAL ONLY — DO NOT SHARE WITH PROSPECTS***

**What they are:** Agentic process automation platform targeting mid-market operations teams. Agent-based workflow automation with pre-built industry templates.

**Their strengths:**

- Strong operations workflow templates for finance, HR, and supply chain
- Clean UI with visual workflow builder
- SOC 2 Type II certified (we are not — yet)
- $12M Series A, established team

**Their weaknesses:**

- Operations-only — no creative, no engineering, no media production
- 4 verticals only (finance, HR, supply chain, customer support)
- AWS-only infrastructure — no multi-cloud
- $1,500 entry point with $750 setup fee
- 7-day pilot only (our Free Assessment + 14-day trial is better)
- No model transparency — usage costs buried in pricing

**Cavi's Defensible Wins Against WorkflowAS:**

| Dimension | WorkflowAS | Cavi |
|---|---|---|
| Pricing tiers | 2 tiers | 9 tiers — covers solo to enterprise |
| Entry point | $1,500/mo + $750 setup | $1,200/mo + $0 setup |
| Free offering | 7-day pilot | Free Workflow Assessment + 14-day free trial (Solo/Duo/Starter) |
| Coverage | 4 verticals (ops only) | 8 divisions spanning ops, creative, engineering, research, media, talent |
| Specialists | ~30 (operations-focused) | 66+ across operational, creative, technical, and strategic — and growing |
| Infrastructure | AWS-only | Multi-cloud (AWS, GCP, Azure) + hybrid deployment |
| Model transparency | Costs buried | Model Budget included, wholesale pass-through, zero markup |

**Talk track (internal preparation — never name them):**

> "*Some platforms focus on automating one type of work — usually operations. We cover eight coordinated divisions because real businesses do not run on operations alone. You need marketing, engineering, research, media, talent, and project management working together. That is what coordination means.*"

---

## 7.3 PaperclipAI — Foundation Architecture

> ***CRITICAL: PaperclipAI is NOT a competitor. It is a foundational technology that Cavi is evaluating and plans to build upon.***

PaperclipAI provides foundational open-source agent architecture that Cavi is exploring as a base layer for future development. Rather than competing with PaperclipAI, our roadmap calls for extending it — layering on proprietary orchestration, division structure, human-in-the-loop governance, and enterprise-grade deployment on top of its proven infrastructure.

**Where Cavi plans to extend the foundation:**

| Layer | PaperclipAI Offers | Cavi Will Add |
|---|---|---|
| Agent primitives | Core agent framework, task execution patterns | Named specialists with narrow scopes and defined capabilities |
| Coordination | Basic inter-agent communication | Coordinated division structure — 8 divisions with orchestrators managing specialist teams |
| Governance | Open-source flexibility | Human approval gates at every output stage — nothing ships without review |
| Deployment | Self-hosted, developer-configured | Enterprise deployment: Kubernetes, multi-cloud, managed service, hybrid options |

| | | |
|---|---|---|
| Business layer | Open-source (free, unsupported) | Pricing tiers, onboarding, account management, SLAs, contractual guarantees |

**As PaperclipAI's framework improves, Cavi stands to benefit directly — we are aligned, not competing.** This is a strength, not a weakness. Building on a strong open-source foundation means the underlying primitives get better without us having to rebuild them.

**If a prospect mentions PaperclipAI:**

> *"We're evaluating PaperclipAI's architecture as a foundation layer — they provide a strong open-source base for multi-agent systems. What we plan to add on top is the orchestration, the division structure, the human approval gates, and the managed deployment that turns a framework into a workforce you can direct without engineering. Think of it like the difference between Linux and a managed cloud service — the foundation is open-source, but the business value is in what you build on top of it."*

**If a technical prospect wants to go deeper:**

> *"We're looking at PaperclipAI for its proven agent primitives. What we layer on top is what makes it operational: first, a coordination layer that organizes specialists into divisions with orchestrators. Second, a governance layer that enforces human approval at every gate. Third, a deployment layer that handles enterprise infrastructure — multi-cloud, Kubernetes, managed service, monitoring, and SLAs. Happy to have our engineering lead walk through the architecture in detail."*

---

## 7.4 AI Employee Category (e.g., Devin, Harvey, Jasper)

**What they are:** Single-purpose AI tools branded as "employees" — an AI developer, an AI lawyer, an AI writer.

**Where they fall short:**

- Single-function only — great at one thing, useless for everything else
- No cross-functional coordination — the AI writer does not talk to the AI developer
- Pricing per function adds up fast — $200/mo for a writer + $400/mo for a coder + $300/mo for a researcher = $900/mo for three disconnected tools
- No human-in-the-loop governance — outputs ship directly without review in many cases

**Cavi's position:**

> *"Those tools are excellent at one thing. The question is: does your business only need one thing? Most operators need marketing, engineering, research, operations, and media working together. Cavi gives you 66+ specialists across 8+ coordinated divisions — all working under your team's direction, all coordinating with each other. One platform, one subscription, one approval workflow."*

---

## 7.5 Enterprise Agent Platforms (e.g., Salesforce Einstein, Microsoft Copilot)

**What they are:** AI capabilities embedded into existing enterprise platforms.

**Where they fall short:**

- Locked into one ecosystem — Salesforce agents only work in Salesforce, Copilot only works in Microsoft 365
- Expensive — enterprise AI add-ons often cost $30-$50/user/month on top of existing licensing
- General-purpose — not specialized for specific workflows

- Slow deployment — months of configuration, training, and change management

**Cavi's position:**

> *"Enterprise AI add-ons are ecosystem-locked. If your whole business runs in one ecosystem, they can help. But most businesses use multiple tools across multiple platforms. Cavi is platform-agnostic — 66+ specialists that integrate with your existing stack, not one vendor's stack. And we deploy in 3 days, not 3 months."*

---

## 7.6 Traditional Agencies and Offshore Teams

**What they are:** Human services firms — marketing agencies, development shops, offshore staffing companies.

**Where they fall short:**

- Slow to start — 2-4 weeks minimum onboarding
- Expensive at scale — a mid-tier agency runs $10K-$50K/month
- Knowledge walks out the door when the engagement ends
- Quality inconsistency — depends on which humans are assigned to your account
- Time zone and communication overhead for offshore

**Cavi's position:**

See Objection #14 in Part 6 for the full talk track. The short version: agencies are great for high-strategy work. Cavi is great for high-volume execution. Many clients will use both.

---

## 7.7 General-Purpose AI Tools (ChatGPT, Claude, Gemini)

**What they are:** General-purpose AI assistants available directly to end users.

**Where they fall short:**

- No specialization — same model tries to do everything
- No coordination — each conversation is isolated
- No workflow integration — copy-paste between tools
- No audit trail — conversations disappear
- No team governance — no approval gates, no RBAC, no accountability

**Cavi's position:**

See Objection #6 in Part 6 for the full talk track. The short version: general AI tools are search engines. Cavi is a research team. Keep the tools for ad hoc questions. Use Cavi for the work.

---



---

# Part 8. Sales Scripts

---

## 8.1 SMB Primary Discovery Call (30 minutes)

**Objective:** Qualify, build trust, present solution, close to pilot or subscription.

---

**Minutes 0-2: Open warm.**

> *"Thanks for making the time. I know your calendar is packed, so I want to make sure this is useful for you — not just a pitch. Can you give me 30 seconds on what prompted you to take this call?"*

*Listen. Take notes. Whatever they say is your anchor for the rest of the call.*

---

**Minutes 2-10: Discovery.**

> *"Tell me about your team. How many people, and what are they spending most of their time on?"*

> *"What is the work that is not getting done right now — the stuff that keeps getting pushed to next week?"*

> *"How are you handling [their specific pain] today? In-house, agency, freelancers, or just not doing it?"*

> *"What have you tried before? What worked and what did not?"*

> *"If nothing changes in the next 90 days, what does that cost you — in time, money, or missed opportunities?"*

---

**Minutes 10-18: Present the solution.**

> *"Here is what I am hearing: [summarize their pain in their words]. Let me show you how Cavi addresses that specifically.*
>
> *Cavi is a managed AI workforce — 66+ specialists organized across 8 coordinated divisions. Each specialist has a narrow job and works under your team's direction. Nothing ships without your approval.*
>
> *For what you described, I would recommend [specific tier or specialist configuration]. Here is what that looks like in practice: [walk through a day-in-the-life scenario using their specific workflow].*
>
> *The timeline is 3 days. Day 1 — onboarding call and configuration. Day 2 — integration and setup. Day 3 — first output reviewed by your team. You would have a working [specialist type] by [specific day of the week]."*

---

**Minutes 18-25: Handle objections and price.**

> *"The [recommended tier] is [$X] per month. That includes [Y specialists], the Model Budget, and $0 setup. Month-to-month — no long-term commitment.*
>
> *To put that in context, one employee doing this work costs [$80-120K/year loaded]. You are getting [Y specialists] for [$X/month] — projected to handle the workload of [Y to 3Y employees] based on our internal usage.*
>
> *What questions do you have about the pricing?"*

*Address objections using Part 6.*

---

**Minutes 25-30: Close.**

> *"Based on everything we discussed, here is what I would recommend: [specific next step].*
>
> ***Option A (if they are ready):*** *Let us start with [tier] today. I can have your first specialist deployed by [day]. I will send the agreement over this afternoon.*
>
> ***Option B (if they need more):*** *Let us do the Free Workflow Assessment. I will map your workflows, identify the highest-leverage specialist, and come back with a specific recommendation. Can we schedule that for [tomorrow/Thursday]?*
>
> ***Option C (if they want proof):*** *The SMB Mini-Pilot is $1,500 for 3 days. We deploy a specialist on your actual workflow, produce real output, and deliver four deliverables. If you move forward, the $1,500 is credited to your first month. Can we kick that off this week?"*

---

## 8.2 Solopreneur Fast-Close Call (15 minutes)

**Objective:** Qualify fast, connect emotionally, close to Solo trial.

---

**Minutes 0-2: Connect.**

> *"I know your time is the most valuable thing you have, so let me keep this tight. Tell me about your business — what you do and where you feel stretched thinnest."*

---

**Minutes 2-7: Discovery (compressed).**

> *"What is the one thing that if it were handled for you, would change your week?"*

> *"How much time are you spending on that right now? Per week, roughly?"*

> *"What have you tried — freelancers, VAs, tools? What happened?"*

> *"If I told you there is a specialist that handles [their pain] 24/7 for $1,200 a month, what would your first reaction be?"*

---

**Minutes 7-12: Present and handle.**

> *"Here is what I would do for you. One specialist — [specific to their need]. Deployed in 3 days. Runs 24/7. Everything it produces goes through you for approval before it goes anywhere.*
>
> *$1,200 a month. That is less than one freelancer invoice. And unlike a freelancer, it does not miss deadlines, take vacation, or need to be re-briefed every time.*
>
> *You get 14 days free — no card required. If it does not change your week, walk away. No friction."*

*Handle any objection from Part 6. Solopreneurs typically raise #5 (price), #9 (too small), or #7 (need to think).*

---

**Minutes 12-15: Close.**

> *"Here is the next step: I set you up with the 14-day trial right now. You will have your specialist by [day]. If it works, you subscribe. If not, you have lost nothing.*
>
> *Can we get started?"*

---

## 8.3 Enterprise Pilot Pitch (45 minutes)

**Objective:** Build credibility, establish executive trust, close to $10K Enterprise Pilot.

---

**Minutes 0-5: Executive alignment.**

> *"Thank you for the time. Before I share anything about Cavi, I want to understand where you are. What is driving your evaluation of AI workforce solutions right now?"*

> *"What does success look like for you specifically — not for the company, for you — by the end of this year?"*

> *"What AI initiatives have you tried before? What worked and what did not?"*

---

**Minutes 5-15: Deep discovery.**

> *"Walk me through the workflows that are creating the most pressure right now. Where is your team spending time on work that does not require human judgment?"*

> *"How many people touch that process? What is the fully loaded cost?"*

> *"What does your security review process look like for new vendors? Who needs to be in the room?"*

> *"What budget is allocated for AI or workforce transformation this year?"*

> *"If we could deploy a pilot in the next 30 days that answers every question your team has — security, performance, ROI — would that be the right next step?"*

---

**Minutes 15-30: Platform presentation.**

> *"Let me show you how Cavi works and then we will map it to your specific workflows.*
>
> *Cavi is the Agentic Operations Platform — a managed AI workforce. 66+ specialists organized across 8 coordinated divisions — and growing. Each specialist has a narrow scope, works under your team's direction, and produces output that passes through human approval gates before it goes anywhere.*
>
> *The divisions are: Executive Operations, Engineering & Development, Marketing & Brand, AI Media Production, Talent Acquisition, Project Operations & Analytics, Research & Intelligence, and Infrastructure & Reliability.*
>
> *For your team, I would recommend starting with [specific divisions based on discovery]. Here is what that looks like."*

*Walk through a specific scenario using their workflows. Be concrete — use their language, their bottlenecks, their numbers.*

> *"On the enterprise side, we provide SSO, RBAC, full audit trails, data isolation, and hybrid deployment options — Kubernetes on your infrastructure or Cavi-hosted. Our engineering lead can walk through the architecture with your security team in a dedicated session."*

---

**Minutes 30-40: Pilot proposal.**

> *"Here is what I recommend: the 30-day Enterprise Pilot.*
>
> *$10,000. Four deliverables:*
>
> 1. ***Workflow Constraint Map** — we map exactly where your current process breaks down and where Cavi adds leverage.*
> 2. ***Risk & Control Design** — how Cavi handles your specific security and compliance requirements.*
> 3. ***Implementation Blueprint** — phased rollout plan with timelines, success metrics, and resource requirements.*
> 4. ***Executive Decision Ready Readout** — board-ready summary with ROI projections and a go/no-go recommendation.*
>
> *At the end of 30 days, you have everything your team needs to make a confident decision. If you proceed, the $10,000 is credited toward your first quarter. If you do not, you keep the Constraint Map and the Blueprint.*
>
> *This is not a demo with a timer. It is a real engagement with real deliverables."*

---

**Minutes 40-45: Close and next steps.**

> *"Who else needs to be in the room for the pilot decision? Let us get a follow-up with your [security lead / CTO / procurement] scheduled this week. I will send the pilot proposal in writing this afternoon."*

---

# Part 9. Outreach Templates

## 9.1 SMB LinkedIn Connection Request

> *Hi [First Name] — I help operators at companies like [Company] add AI specialists that handle high-volume work without adding headcount. We have 66+ specialists across 8 divisions deployed in 3 days. Worth a 10-minute conversation?*

---

## 9.2 SMB Cold Email — First Touch

**Subject:** [Company] + Cavi — capacity without headcount

> *Hi [First Name],*
>
> *I noticed [specific observation — recent hire posting, product launch, funding round, etc.]. Congratulations.*
>
> *Quick question: is your team capacity-constrained right now? Most operators at your stage are — more demand than people to deliver against it.*
>
> *Cavi is a managed AI workforce — 66+ specialists organized across 8 coordinated divisions. We deploy in 3 days, your team reviews and approves everything, and it costs a fraction of one hire.*
>
> *Would a 10-minute call this week make sense? I can show you exactly where AI specialists would fit in your workflows.*
>
> *Best, [Your Name] [Title], Cavi Collective www.cavivaultagents.io*

---

## 9.3 SMB Cold Email — Second Touch (Day 3)

**Subject:** Re: [Company] + Cavi — capacity without headcount

> *Hi [First Name],*
>
> *Following up on my note from [day]. I know inboxes are brutal.*
>
> *One thing I should have led with: we offer a Free Workflow Assessment — 30 minutes, no pitch, no commitment. I map your current workflows, identify the bottlenecks, and show you where AI specialists would add the most leverage.*
>
> *Even if Cavi is not the right fit, you walk away with a workflow map. Worth the 30 minutes?*
>
> *[Your Name]*

---

## 9.4 SMB Cold Email — Third Touch (Day 7)

**Subject:** Last note — [First Name]

> *Hi [First Name],*
>
> *Last note from me. If the timing is not right, I completely understand.*

> *If capacity becomes a pressure point in the next quarter — and for most growing companies it does — here is what we do: 66+ AI specialists across 8 divisions, deployed in 3 days, working under your team's direction.*
>
> *I am here when the timing is right.*
>
> *[Your Name] www.cavivaultagents.io*

---

## 9.5 Solopreneur LinkedIn Connection Request

> *Hi [First Name] — I saw your work on [specific project/post]. Impressive what you are building solo. Quick question: what would change if you had a dedicated specialist handling [their likely pain — content, research, operations] 24/7 for less than a freelancer invoice?*

---

## 9.6 Solopreneur Cold Email

**Subject:** A specialist for [their business type] — $1,200/mo, live in 3 days

> *Hi [First Name],*
>
> *Building [their type of business] alone is exhausting. Every hour you spend on [ops/content/admin] is an hour not spent on revenue.*
>
> *Cavi gives you a dedicated AI specialist — deployed in 3 days, runs 24/7, and everything goes through you for approval. $1,200 per month. Less than one freelancer project.*
>
> *14 days free. No card. If it does not change your week, walk away.*
>
> *Want to try it?*
>
> *[Your Name] www.cavivaultagents.io*

---

## 9.7 Enterprise Cold Email — First Touch

**Subject:** [Company] AI workforce pilot — 30 days, 4 deliverables

> *Hi [First Name],*
>
> *I work with operations leaders at companies like [Company] who need to show measurable AI ROI — not another pilot that dies in committee.*
>
> *Cavi is a managed AI workforce — 66+ specialists across 8 coordinated divisions. We run a 30-day Enterprise Pilot that delivers four things:*
>
> 1. *Workflow Constraint Map — where your current process breaks down*
> 2. *Risk & Control Design — how we handle your security requirements*
> 3. *Implementation Blueprint — phased rollout with timelines*
> 4. *Executive Decision Ready Readout — board-ready summary*
>
> *$10,000. Credited to your first quarter if you proceed. You keep the deliverables either way.*
>
> *Is this worth a 20-minute conversation this week?*
>
> *[Your Name] [Title], Cavi Collective www.cavivaultagents.io*

---

## 9.8 Enterprise Cold Email — Second Touch (Day 5)

**Subject:** Re: [Company] AI workforce pilot

> *Hi [First Name],*
>
> *Following up. One thing that often resonates with leaders in your position: most AI pilots fail because they are science projects with no defined deliverables.*
>
> *Our Enterprise Pilot is the opposite — fixed scope, fixed timeline, four deliverables, and a go/no-go recommendation your team can take to the board.*
>
> *20 minutes to see if it fits?*
>
> *[Your Name]*

---

## 9.9 Follow-Up After Free Workflow Assessment

**Subject:** Your Workflow Assessment — next steps

> *Hi [First Name],*
>
> *Thanks for the conversation today. Attached is your Mini Workflow Map.*
>
> *Based on what we mapped, here is my recommendation:*
>
> *[Specific tier and specialist configuration] [Monthly price] [What the first deliverable would be] [Timeline to deploy]*
>
> *The next step is [specific: start trial / schedule decision call / kick off pilot]. Can we lock that in for [specific day]?*
>
> *[Your Name]*

---

## 9.10 Re-Engagement Email (30+ days since last contact)

**Subject:** Still relevant, [First Name]?

> *Hi [First Name],*
>
> *We spoke [X weeks/months] ago about [specific topic from the conversation]. You mentioned [specific pain point].*
>
> *Wanted to check in: is that still a pressure point? We have added new capabilities since we last spoke — 66+ specialists across 8 divisions now, with [mention any relevant new division or capability].*
>
> *If the timing is better now, I would love to pick up where we left off. Free Workflow Assessment is still on the table — 30 minutes, no commitment.*
>
> *[Your Name] [www.cavivaultagents.io](http://www.cavivaultagents.io)*

---

## 9.11 Founding Customer Pitch Email

**Subject:** Founding customer invitation — [Company] + Cavi

> *Hi [First Name],*
>
> *I am reaching out with something we do not offer publicly: a founding customer invitation.*

> *Cavi is a managed AI workforce — 66+ specialists across 8 coordinated divisions — launching with a small group of hand-selected founding customers. Founding customers get three things that disappear after launch:*
>
> 1. ***Founder-level attention.*** *The founders are in the room — on your calls, shaping your deployment, personally accountable.*
> 2. ***Founding pricing protections.*** *The price you lock in now is protected for the life of your subscription. As we grow and prices adjust, yours does not.*
> 3. ***Product influence.*** *A real seat at the table on what we build next. Your feedback directly shapes the roadmap.*
>
> *We are looking for operators who want to shape the product, not just use it. Based on [specific reason — their company, industry, use case], I think [Company] would be an ideal founding partner.*
>
> *Can we schedule 20 minutes this week?*
>
> *[Your Name] [Title], Cavi Collective [www.cavivaultagents.io](http://www.cavivaultagents.io)*

---



---

# Part 10. Proof Points & Case Study

---

> ***PRE-LAUNCH HONESTY RULES***
>
> *We are pre-launch. We do not have external client results. Every claim in this section must be framed as a projection from internal usage. Violating this rule destroys trust and creates legal exposure. There are no exceptions.*
>
> ***Approved language:*** *"We project..." / "Based on internal usage..." / "Working projection, not a customer case study..."*
>
> ***Forbidden language:*** *"Our clients see..." / "Customers report..." / "Case study results show..."*

---

## 10.1 Approved Projections (Memorize Both)

**Projection 1: 8-10x Output Capacity Gain**

> *"Based on Cavi's partner running the platform internally, we project an 8 to 10 times output capacity gain when a coordinated AI team is paired with a human operator. That is not a marketing number — it is a working projection from real usage. We share it as a projection because we are early and we would rather you trust us in two years than oversell you today."*

**Projection 2: 1 Specialist = 1-3 Employees (Projected)**

> *"Based on Cavi's partner running the platform internally, we project a single specialist can handle the workload of 1 to 3 employees depending on the workflow. A specialist runs continuously, handles multi-step workflows in parallel, and does not take vacation. Same honesty framing — projection from internal usage, not a customer case study."*

**How to deliver projections:**

1. State the number clearly and confidently.
2. Immediately frame it as a projection from internal usage.
3. Acknowledge pre-launch status without apologizing for it.
4. Pivot to the pilot: "That is why the pilot exists — so you can see your own numbers on your own workflows."

---

## 10.2 Internal Case Study — Cavi Running Cavi

**Context:** Cavi uses its own platform internally to run its operations. This is the only case study we have, and it is powerful — but it must be framed correctly.

**The story:**

> *"We run Cavi on Cavi. Our marketing, engineering, research, media production, and project operations all run through the platform. A small founding team operates with the output capacity of a much larger organization because 66+ specialists handle the execution layer while the human team focuses on strategy, quality, and direction.*
>
> *We built the platform because we needed it ourselves. That is why it works — it was not built in a lab. It was built in the field, by operators, for operators."*

**The numbers (internal, projected):**

- Content output: projected 8-10x increase in volume with same human team
- Engineering velocity: projected 3-5x increase in features shipped per sprint
- Research capacity: projected 4-6x increase in reports and analyses produced
- Media production: projected 5-8x increase in assets created per week

> ***Framing for all numbers:*** *"These are projections from our own internal usage. We are not claiming they will be identical for every business. But they give you a sense of the capacity expansion that is possible when specialists and humans work together."*

---

## 10.3 Founding Customer Frame

When a prospect asks for references, customer stories, or case studies:

> *"We are pre-launch and hand-selecting our founding customer group. That means we do not have external case studies yet — and we are honest about that. What we can offer is three things:*
>
> *First, the internal case study: we run the platform ourselves and can show you the output in real time.*
>
> *Second, the pilot: 3 days or 30 days depending on your needs. Real output on your workflows. Your own numbers.*
>
> *Third, the founding customer relationship: you get founder-level attention, founding pricing protections, and direct product influence. Our first customers do not just use the product — they shape it."*

---

## 10.4 Hypothetical Scenario Template

Use these to help prospects visualize Cavi in their specific context. Always frame as hypothetical — "here is how this would work for a company like yours."

**Template — Marketing Team:**

> *"Imagine your marketing team today produces 10 pieces of content per week across blog, social, and email. With the Marketing & Brand division — 25 specialists across 5 departments — we project that same team could produce 80-100 pieces per week, all reviewed and approved by your team before publishing. The human team focuses on strategy, voice, and quality. The specialists handle research, drafting, formatting, scheduling, and analytics. The result: 8-10x more output with the same team."*

**Template — Engineering Team:**

> *"Your engineering team ships features on a 2-week sprint cycle. With the Engineering & Development division — 15 specialists — we project that cycle compresses to a few days for many feature types. The specialists handle boilerplate, testing, documentation, code review prep, and deployment automation. Your engineers focus on architecture decisions, complex logic, and code review. The result: more features shipped, less burnout."*

**Template — Solo Operator:**

> *"You are running everything alone — content, outreach, admin, research, client delivery. One specialist at $1,200 per month handles [their specific bottleneck]. You get back 10-15 hours per week. That is 10-15 hours redirected to revenue-generating work. At your billing rate of $[their rate], that is $[math] per month in recaptured revenue capacity — from a $1,200 investment."*

---

## 10.5 References Response

When a prospect asks to speak with a current customer:

> *"We are pre-launch, so we do not have external references in the traditional sense. Here is what I can offer instead:*
>
> 1. ***A live demo of the platform running our own operations.*** *You can see real output, real workflows, real coordination between specialists.*
> 2. ***The pilot.*** *Instead of hearing someone else's experience, you get your own. Real output on your workflows, in 3 or 30 days.*
> 3. ***A conversation with the Founders.*** *They can speak to the platform's capabilities, the architecture, the roadmap, and the honest state of the business.*
>
> *As we onboard founding customers, we will have references. But right now, the most credible evidence is the pilot — your workflows, your output, your numbers."*

---



---

# Part 11. Technical Breakdown Bonus

---

> ***THE RULE: Share categories, never implementations.***
>
> *When a prospect asks a technical question, answer at the category level. "We use containerized microservices orchestrated through Kubernetes on multi-cloud infrastructure." Do NOT say which specific cloud provider, which Kubernetes distribution, which specific model version, or which internal tool names. Categories are public. Implementations are proprietary.*
>
> *If the prospect needs implementation-level detail for a security review, escalate to the engineering lead. Do not guess.*

---

## 11.1 Why This Matters

Technical credibility closes enterprise deals. An Account Executive who can answer the first two technical questions confidently and escalate the third one correctly is worth more than one who hands off every technical question to an engineer.

**Your goal in technical conversations:**

1. Answer confidently at the category level.

2. Demonstrate that the platform is real, production-grade, and built by experienced engineers.
3. Escalate gracefully when the prospect goes deeper than categories.
4. Never guess. Never bluff. Never make up a technical detail.

---

## 11.2 Foundation Model Glossary

Know these terms. You do not need to explain them in detail — you need to use them correctly.

| Term | Definition | When You Use It |
|---|---|---|
| **Foundation Model** | A large AI model (GPT-4, Claude, Gemini, etc.) trained on broad data. The raw intelligence layer. | "We are model-agnostic — we route between multiple foundation models." |
| **Fine-Tuning** | Training a foundation model on specific data to improve it for a narrow task. | "Our specialists use fine-tuned models where it improves output quality for specific workflows." |
| **Prompt Engineering** | Crafting the instructions given to a model to get better output. | "Each specialist has purpose-built prompts engineered for their specific task." |
| **RAG (Retrieval-Augmented Generation)** | Giving a model access to external data (documents, databases) so it can reference real information instead of relying only on training data. | "Our specialists use RAG to work with your business documents — they reference your actual data, not just their training." |
| **Embedding** | Converting text into a numerical representation that captures meaning. Used for search, similarity, and retrieval. | "We use embeddings to help specialists find relevant information in your documents." |
| **Token** | The basic unit of text that models process. Roughly 0.75 words per token. | "Tokens are how model usage is measured. Our Model Budget covers typical volume." |
| **Context Window** | The amount of text a model can consider at once. Measured in tokens. | "Modern models have large context windows — they can work with long documents in a single pass." |
| **Inference** | The process of a model generating output from input. Each API call is an inference. | "Inference costs are what the Model Budget covers." |
| **Latency** | The time between sending a request to a model and getting a response. | "We optimize for low latency — specialists respond quickly, not in minutes." |

---

## 11.3 Architecture Glossary

| Term | Definition | When You Use It |
|---|---|---|
| **Microservices** | An architecture where each component runs independently and communicates through APIs. | "Each specialist runs as an independent microservice — |

| | | |
|---|---|---|
| | Each specialist is a microservice. | isolated, scalable, and independently deployable." |
| **Kubernetes (K8s)** | An orchestration platform that manages containerized applications at scale. | "We run on Kubernetes — industry-standard orchestration for containerized workloads." |
| **Container** | A lightweight, isolated environment that packages an application and its dependencies. | "Each specialist runs in its own container — fully isolated from other specialists and other clients." |
| **Multi-Cloud** | Running infrastructure across multiple cloud providers (AWS, GCP, Azure) for redundancy and flexibility. | "We are multi-cloud — not locked into any single provider." |
| **API Gateway** | A single entry point that routes requests to the appropriate backend service. | "All requests flow through our API gateway — centralized routing, authentication, and rate limiting." |
| **Message Queue** | A system for asynchronous communication between services. | "Specialists communicate through message queues — reliable, ordered, and auditable." |
| **Load Balancer** | Distributes incoming requests across multiple instances of a service. | "Load balancing ensures consistent performance even under heavy usage." |
| **Data Isolation** | Ensuring one client's data is completely separate from another's. | "Full data isolation — your data never comingles with another client's data." |
| **Division** | A self-contained department within the Cavi platform, led by an orchestrator and staffed with domain specialists. Each division has its own workspace, capability set, and integration stack. | "Cavi is organized into 8 divisions, each led by an orchestrator who coordinates the specialists within it." |
| **Orchestrator** | The lead agent in a division. It delegates work to specialists, coordinates handoffs, and reports back to the operator for approval. Orchestrators do not execute tasks directly — they coordinate. | "Each division has an orchestrator that manages the specialists, coordinates work, and ensures nothing falls through the cracks." |

---

## 11.4 AI Engineering Glossary

| Term | Definition | When You Use It |
|---|---|---|
| **Agent** | An AI system that can take actions, use tools, and work toward goals — not just generate text. | "Our specialists are agents — they take actions, use tools, and produce deliverables." |
| **Tool Use** | An agent's ability to call external tools (APIs, databases, file systems) as part of its work. | "Specialists use tools — they can search databases, call APIs, generate files, and interact with your tech stack." |

| | | |
|---|---|---|
| **Chain of Thought** | A technique where the model reasons step-by-step before producing a final answer. Improves accuracy. | "Our specialists use chain-of-thought reasoning for complex tasks — step-by-step processing, not one-shot guessing." |
| **Model Routing** | Selecting the optimal model for each specific task based on requirements (speed, accuracy, cost, capability). | "We route between multiple models — each task goes to the model that handles it best." |
| **Fallback** | An alternative model or process that activates if the primary one fails or underperforms. | "Every specialist has fallback models — if the primary model is unavailable, the specialist switches automatically." |
| **Guardrails** | Rules and constraints that prevent a model from producing harmful, off-topic, or incorrect output. | "Every specialist has guardrails — scope limits, content filters, and output validation." |
| **Evaluation (Evals)** | Systematic testing of model output quality against defined criteria. | "We run continuous evaluations to ensure specialist output meets quality standards." |

---

## 11.5 Security and Compliance Glossary

| Term | Definition | When You Use It |
|---|---|---|
| **SOC 2 Type II** | An audit standard that verifies a company's security controls over time. The gold standard for SaaS security. | "We are building toward SOC 2 Type II. Here is what we do today..." (never claim we have it) |
| **RBAC (Role-Based Access Control)** | Limiting access based on a user's role — admins see everything, operators see their division, viewers see output only. | "Full RBAC — your team members see only what their role allows." |
| **SSO (Single Sign-On)** | Allowing users to log in with their existing company credentials (Google Workspace, Okta, Azure AD). | "SSO integration — your team logs in with their existing identity provider." |
| **Encryption at Rest** | Data is encrypted when stored on disk. | "All data encrypted at rest using industry-standard encryption." |
| **Encryption in Transit** | Data is encrypted when moving between systems (TLS/HTTPS). | "All data encrypted in transit — TLS everywhere." |
| **Audit Trail** | A complete, tamper-proof log of every action, input, output, and decision. | "Full audit trail — every action logged and reviewable." |
| **Data Residency** | Where data is physically stored (which country, which region). | "We can discuss data residency requirements during scoping." |
| **Penetration Testing** | Security testing where experts attempt to break into the system to find vulnerabilities. | "Penetration testing is part of our security program." |
| **Zero Trust** | A security model that verifies every request, even from inside the network. | "Zero trust architecture — every request is authenticated and authorized." |

## 11.6 Sample Technical Questions and Responses

**Question 1: "Walk me through the architecture."**

> *"Cavi is cloud-native, containerized, and orchestrated through Kubernetes — multi-cloud, not locked into any single provider. 66+ specialists organized as microservices coordinating through our orchestration layer with human approval gates at every output stage.*
>
> *Each specialist runs in its own container — fully isolated. Specialists within a division coordinate through an orchestrator that manages task delegation, handoffs, and quality checks. Divisions coordinate across boundaries when work requires cross-functional execution.*
>
> *Model-agnostic — we route between multiple foundation model providers, selecting the best model for each specific task with automatic fallbacks.*
>
> *For deeper architectural detail, our engineering lead can walk your team through it in a dedicated session. They can go as deep as your team wants."*

**Question 2: "How do specialists work with our business documents?"**

> *"Through RAG — Retrieval-Augmented Generation. Your documents are ingested, chunked, and embedded into a vector store. When a specialist needs information from your documents, it retrieves the relevant sections and uses them as context for its work.*
>
> *This means the specialist references your actual data — your style guides, your product specs, your customer data — not just its training. The documents stay in your isolated environment and are never used to train models.*
>
> *The result: output that is grounded in your business reality, not generic AI text."*

**Question 3: "Which models do you use?"**

> *"We are model-agnostic. We route between multiple foundation model providers — selecting the optimal model for each task type based on accuracy, speed, cost, and capability. When a better model becomes available from any provider, we evaluate and integrate it.*
>
> *This protects you from provider lock-in and ensures you always get the best available model for each type of work. If one provider has a strong reasoning model and another has a strong creative model, your specialists use both.*
>
> *We do not name specific models in sales conversations because the model landscape changes rapidly and we want to give you the best option at any given time, not lock you into a model that was best six months ago."*

**Question 4: "How do you handle multi-tenant data isolation?"**

> *"Complete isolation. Each client's data lives in its own environment — separate storage, separate processing, separate access controls. One client's data never touches another client's data, models, or processing pipeline.*
>
> *Access is controlled through RBAC — role-based access control. Your team members see only what their role allows. Admin roles can see and manage everything. Operator roles see their divisions. Viewer roles see output only.*
>
> *All data is encrypted at rest and in transit. Audit trails capture every access event. For clients with data residency requirements, we discuss geographic constraints during scoping."*

**Question 5: "Is there an audit trail?"**

> *"Yes. Comprehensive. Every input, every output, every decision, every tool call, every model invocation — logged, timestamped, and attributable. The audit trail is tamper-proof and exportable.*
>
> *Your team can review any specialist's work history, see exactly what inputs produced what outputs, and trace any decision back to its source. This is not optional — it is built into the platform. If something goes wrong, you can trace it. If an auditor asks, you can show them."*

---

**Question 6: "Does our data get used to train AI models?"**

> *"No. Never. Your data is used to serve your specialists — through RAG, through context, through your specific configurations. It is never used to train or fine-tune any foundation model. It is never shared with model providers. It is never used to improve models for other clients.*
>
> *This is contractual, not just policy. Your MSA explicitly prohibits the use of your data for model training. Your data is yours. Period."*

---

## 11.7 When to Bring In Technical

**Escalation triggers — bring in the engineering lead when:**

- The prospect asks about specific encryption algorithms, key management, or cryptographic standards
- The prospect asks about specific compliance certifications (SOC 2 audit details, HIPAA BAA terms, FedRAMP authorization)
- The prospect asks about penetration testing results, vulnerability reports, or security incident response procedures
- The prospect asks about specific model versions, model fine-tuning details, or prompt engineering approaches
- The prospect asks about data residency in specific jurisdictions with legal requirements
- The prospect's security team wants a technical deep-dive or architecture review
- You are not confident in your answer — escalate rather than guess

**Escalation script:**

> "*That is a great question and the right person to walk you through it is our engineering lead. They can go as deep as your team wants on the architecture, the security controls, and the implementation details. Can I set up a 30-minute technical session with them this week?*"

**How to escalate well:**

1. Acknowledge the question: "Great question."
2. Give the category-level answer if you can: "At a high level, we do X."
3. Escalate for depth: "For the implementation specifics, our engineering lead is the right person."
4. Offer a specific next step: "Can I schedule a 30-minute technical session?"
5. Brief the engineering lead before the call with the specific questions.

---

## 11.8 What to Never Say in Technical Conversations

| Never Say This | Why | Say This Instead |
|---|---|---|
| Specific model names and versions to clients | Models change; commitments to specific models create problems | "We route between multiple foundation models, selecting the best for each task." |

| | | |
|---|---|---|
| "Our models do not hallucinate" | All models can hallucinate; lying about it destroys credibility | "We have four layers of defense against hallucination: narrow scope, human approval, audit trails, and model routing." |
| Specific cloud provider names in client-facing materials | Limits flexibility and creates vendor-lock-in perception | "Multi-cloud infrastructure with redundancy." |
| Any compliance certification we do not have | Legal and regulatory exposure | "We are building toward [certification]. Here is what we do today." |
| Specific performance benchmarks we have not verified | Creates commitments we may not meet | "We project [X] based on internal usage. The pilot will show your specific numbers." |
| Internal tool names or proprietary system names | Proprietary information | Use category language: "our orchestration layer," "our routing system." |
| "We are more secure than [competitor]" | Unverifiable and unprofessional | "Here is what our security architecture looks like." |
| "We will have [feature] by [date]" unless confirmed by the engineering lead | Roadmap commitments in sales conversations create delivery pressure and broken promises | "That is on our roadmap. I can get you a timeline from the development team." |



## End of Playbook

**CAVI COLLECTIVE** Sales Enablement Playbook v2.0 Internal & Confidential

This document is the property of Cavi Collective and is intended for internal use only. Distribution outside of authorized personnel is prohibited.

**www.cavivaultagents.io**

*66+ agents across 8 divisions — and growing.*